

# **B565-Data Mining**

## **Homework #6**

Due on Friday, March 31, 2023, 08:00 p.m.

*Instructor: Dr. H. Kurban, Head TA: Md R. Kabir*

**Prem Amal**

March 31, 2023

## Expectation-Maximization Algorithm

This part is provided to help you implement the expectation-maximization algorithm.

```

1: function EXPECTATION-MAXIMIZATION(D,  $k$ ,  $\epsilon$ )
2:    $t \leftarrow 0$ 
3:   Randomly initialize  $\mu_1^t, \dots, \mu_k^t$  ▷ Initialization
4:    $\Sigma_i^t \leftarrow \mathbf{I}, \forall i = 1, \dots, k$ 
5:    $P^t(C_i) \leftarrow \frac{1}{k}, \forall i = 1, \dots, k$ 
6:   repeat
7:      $t \leftarrow t + 1$ 
8:     for  $i = 1, \dots, k$  and  $j = 1, \dots, n$  do ▷ Expectation step
9:        $w_{ij} \leftarrow \frac{f(\mathbf{x}_j | \mu_i, \Sigma_i) \cdot P(C_i)}{\sum_{a=1}^k f(\mathbf{x}_j | \mu_a, \Sigma_a) \cdot P(C_a)}$  ▷ posterior probability  $P^t(C_i | \mathbf{x}_j)$ 
10:    end for
11:    for  $i = 1, \dots, k$  do ▷ Maximization Step
12:       $\mu_i^t \leftarrow \frac{\sum_{j=1}^n w_{ij} \mathbf{x}_j}{\sum_{j=1}^n w_{ij}}$  ▷ re-estimate mean
13:       $\Sigma_i^t \leftarrow \frac{\sum_{j=1}^n w_{ij} (\mathbf{x}_j - \mu_j)(\mathbf{x}_j - \mu_j)^T}{\sum_{j=1}^n w_{ij}}$  ▷ re-estimate covariance matrix
14:       $P^t(C_i) \leftarrow \frac{\sum_{j=1}^n w_{ij}}{n}$  ▷ re-estimate priors
15:    end for
16:  until  $\sum_{i=1}^k \|\mu_i^t - \mu_i^{t-1}\|^2 \leq \epsilon$ 
17: end function

```

## Problem 1

Implement Expectation-Maximization (EM) algorithm for Gaussian mixture models (see the EM algorithm above) in *R* or *Python* and call this program  $G_k$ . As you present your code explain your protocol for [20 points]

1. initializing each Gaussian
2. maintaining  $k$  Gaussian
3. deciding ties
4. stopping criteria

## R/Python Code

subsectionR/Python script

```
# Sample R Script With Highlighting
```

```

# Sample Python Script With Highlighting
import pandas as pd
import swifter
import matplotlib.pyplot as plt
5 import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
from scipy.stats import multivariate_normal
10 from sklearn.metrics import silhouette_score, calinski_harabasz_score

```

```

from scipy.spatial.distance import euclidean

df= pd.read_csv('dataset_diabetes/diabetic_data.csv')

15 from sklearn.metrics import silhouette_score,
   calinski_harabasz_score
def GMM_initialization(df,k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
20     means_matrix = df.sample(n=k).values
    identity_matrix=np.eye(number_of_columns)
    covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)
    return
25     means_matrix,covariance_matrix,weights_matrix,number_o
       f_rows,number_of_columns

def
30 calculate_posterior(data,means_matrix,covariance_matrix,we
   ights_matrix,k,number_of_columns):
    posterior=np.zeros(k)
    for i in range(k):
        try:
35             pseudo_inverse =
                np.linalg.pinv(covariance_matrix[i] +
                np.eye(covariance_matrix[i].shape[0]) * 1e-2,
                rcond=1e-10)
                posterior[i] = multivariate_normal.pdf(data,
40                 mean=means_matrix[i], cov=pseudo_inverse)
        except Exception as e:
            continue
    return
    posterior*weights_matrix/(posterior*weights_matrix).su
45     m()

def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
    return_counts=True)
50     missing_labels=[i for i in range(k) if i not in
        unique_values]
    unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

    indices=[i for i in range(len(labels)) if
55     labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
    size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
        labels[val]=missing_labels[i]
60     return labels

def sum_of_square_error_em(new_centroids, data, labels):

```

```

columns = data.columns
65 # Join the data dataframe and the labels dataframe
data = data.join(labels)
# Rename the '0' column of the labels dataframe to
'Label'
data.rename(columns={0:'Label'}, inplace=True)
70 sse = []
# Compute the distance between each data point and
its assigned centroid
for i in range(len(new_centroids)):
    distance = np.sum(np.square(data[data['Label']==i]
75 [columns] - new_centroids.iloc[i]
[columns],dtype=np.float64), axis=1)
    #print(distance)
    sse.append(distance.sum())
# Return the sum of squared errors
80
a=sum(sse)
return a

def Calinski_index_em(df_data,clusters):
85 ch_score = calinski_harabasz_score(df_data, clusters)
return ch_score

90 def GMM(df_cleaned_dia,k,tao):
    scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

95 scaled_input_df=
pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

means_matrix,covariance_matrix,weights_matrix,number_o
100 f_rows,number_of_columns=
GMM_initialization(scaled_input_df,k)
likelihood=0
means_matrix_initial=means_matrix
for i in range(k):
105     try:

        pseudo_inverse =
        np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0])
        * 1e-10))
110        likelihood=likelihood+weights_matrix[i]*multiv
        ariate_normal.logpdf(scaled_input,means_matrix
        [i], pseudo_inverse)
    except Exception as e:
        continue
115 log_likelihood_old=np.sum(likelihood)
old_means_matrix_df=pd.DataFrame(means_matrix)

```

```

posterior_probability = np.zeros((scaled_input.shape[0], k))
iterations=0

120 while (True):
    iterations+=1
    # Expectation
    for i in range(scaled_input.shape[0]):
        posterior_probability[i] =
125         calculate_posterior(scaled_input[i],
            means_matrix, covariance_matrix, weights_matrix,
            k, number_of_columns)

    # Maximization
130 posterior_probability=np.nan_to_num(posterior_prob
    ability, nan=0)
    for i in range(k):
        # Calculating weight
        weight = posterior_probability[:, i].sum()
135         #print(weight)
        # Updating each centroid
        means_matrix[i] = (posterior_probability[:,
            i] @ scaled_input) / weight
        #print(1,means_matrix[i])
140         # Subtracting the mean value from data
        scaled_input_diff = scaled_input - means_matrix[i]

        # Update the covariance matrix
        covariance_matrix[i] =
145         (posterior_probability[:, i] *
            scaled_input_diff.T @ scaled_input_diff) /
            weight

        # Update the weights matrix
150         weights_matrix[i] = weight / number_of_rows

    likelihood=0
    for i in range(k):
155         try:
            pseudo_inverse =
                np.linalg.pinv(covariance_matrix[i] +
                np.diag(np.ones(covariance_matrix[i].shape
                [0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*mu
160             ltivariate_normal.logpdf(scaled_input,mean
                s_matrix[i], sudo_inverse)
        except Exception as e:
            continue
165 log_likelihood_new =np.sum(likelihood)

    new_means_matrix_df=pd.DataFrame(means_matrix)
    distance = []
    for col in new_means_matrix_df.columns:

```

```

170         col_distance =
            euclidean(old_means_matrix_df[col],
                new_means_matrix_df[col])
            distance.append(col_distance)
        tao_calculated=sum(distance)/k

175

        if tao_calculated<
            tao:#log_likelihood_new>log_likelihood_old and
180            100*((log_likelihood_new - log_likelihood_old) /
                log_likelihood_old)<tao:

                print ("Converged")
                labels=np.argmax(posterior_probability,axis=1)
185                labels=maintain_k_clusters(labels,k)
                labels_df=pd.DataFrame(labels)
                means_matrix_df=pd.DataFrame(means_matrix,columns=
                    scaled_input_df.columns)
                sse=sum_of_square_error_em(means_matrix_df,
190                scaled_input_df, labels_df)
                clainski=
                    Calinski_index_em(scaled_input_df,labels_df)
                return sse,clainski,means_matrix_initial
        #else:
195            #log_likelihood_old=log_likelihood_new

        if iterations>100:
            print ("Max iteration reached")
            labels=np.argmax(posterior_probability,axis=1)
200            labels=maintain_k_clusters(labels,k)
            labels_df=pd.DataFrame(labels)
            means_matrix_df=pd.DataFrame(means_matrix,columns=
                scaled_input_df.columns)
            sse=sum_of_square_error_em(means_matrix_df,
205            scaled_input_df, labels_df)
            clainski=
                Calinski_index_em(scaled_input_df,labels_df)
            return sse,clainski,means_matrix_initial

210 error_matrix_em=[]
error_matrix_kmeans=[]
for i in range(2,6):
    for j in range(1,21):
        sse,clainski,means_matrix_initial=GMM(df_cleaned_data,i,10)
215        error_matrix_em.append([i,sse,clainski])

        sse,clainski=kmeans_lyod_with_error(scaled_input_df,i,10,means_matrix_initial)
        error_matrix_kmeans.append([i,sse,clainski])
error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse', 'clainski'])
220 error_df_kmeans= pd.DataFrame(error_matrix_kmeans,columns=['number_of_cluster', 'sse', 'clainski'])

```

```
import seaborn as sns

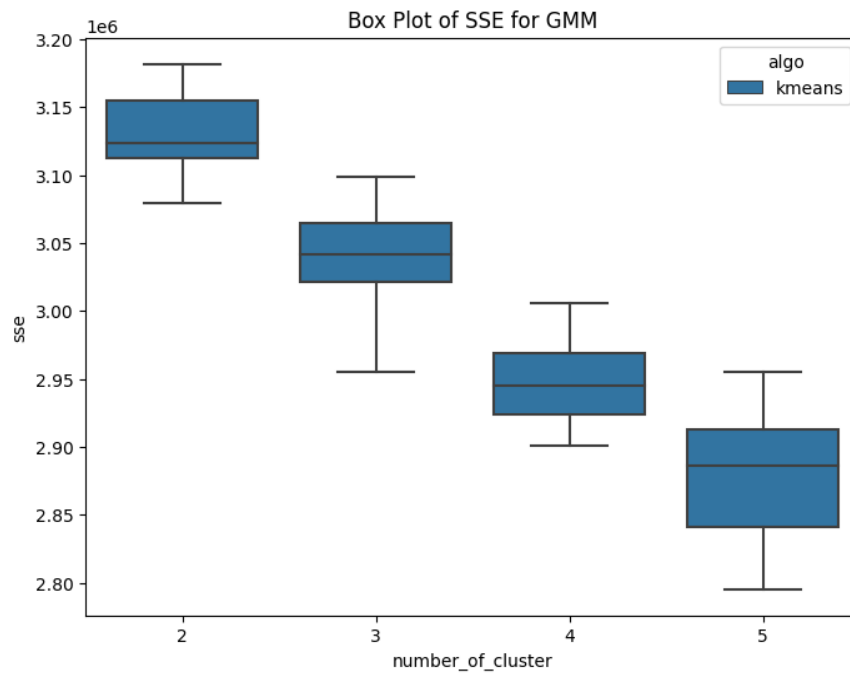
225 fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin
230             (['kmeans'])],ax=ax);
plt.title('Box Plot of SSE for GMM')
plt.show()

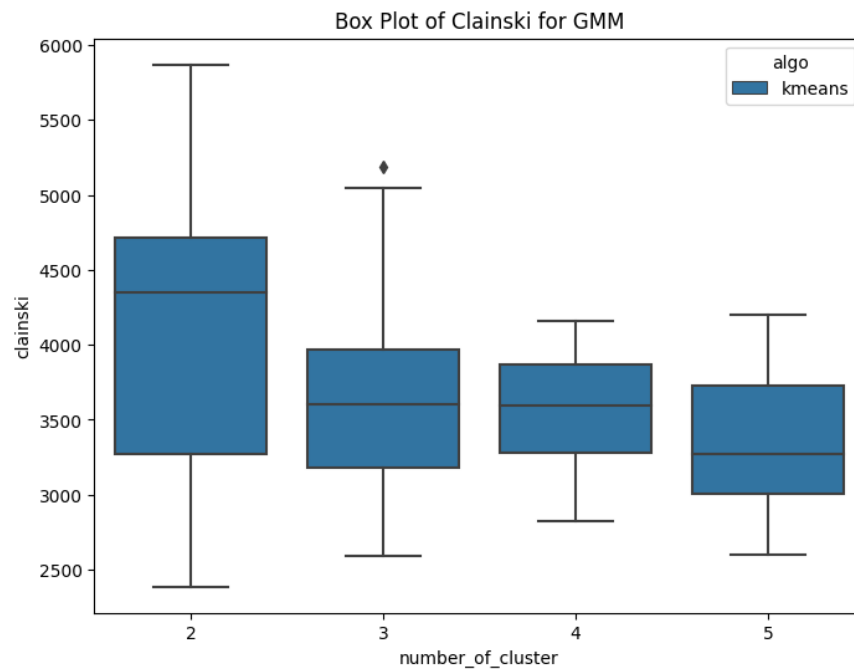
import seaborn as sns

235 fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin (['kmeans'])],ax=ax);
240 plt.title('Box Plot of Clainski for GMM')
plt.show()
```



(1)



(2)

The dataset contains information on patients with diabetes who were admitted to 130 hospitals in the United States between 1999 and 2008. The primary purpose of the dataset is to explore features that affect readmission rates for diabetic patients. The dataset can be used for various purposes, including exploring the relationship between various features like patient condition, labtest, medications and readmission rates, developing predictive models for readmission, and evaluating the quality of care provided by hospitals. The dataset can be used to identify the relationship between various parameters and can be used by Doctors,



researchers to identify the mistakes or areas of improvements or what better could have been done, so that patient is not readmitted. Also, identify some pattern or cluster in the data given all the features or identify the main features affecting the readmission. The different data types that the dataset includes are integer and object data. The dataset's features include patient demographics, admission type and source, diagnosis and procedure codes, lab tests, medications, length of stay, etc. The target variable is whether or not the patient was readmitted to the hospital within 30 days of discharge. Overall there are 101766 entries of data and 50 features/columns including the target variable. There are no null entries as such in data. But some of the columns contains ? Replacing question mark with NAN. After doing that we can see the results and now this missing data needs to be handled. race, weight, payercode, medicalspecialty, diag1, diag2, diag3 Columns contains null data. Dropping the columns where the count of null data is almost around 50 percent as they don't provide significant information.

## Discussion of Initialization of Gaussians

Answer here...

GMM (Gaussian Mixture Model) is a parametric technique. The model is trained using the Expectation-Maximization (EM) algorithm. Before the expectation steps, there are certain assumptions about the data and initializations done, before the start of the actual algorithm. The main assumption is that for all the features, the data is distributed normally and then we use soft clustering to assign each data point to a cluster. The goal of GMM is to estimate the parameters of the underlying mixture model that best fit the observed data. The different vectors initialized are as follows:

- 1) Mean Vector- The mean vector is initialized by randomly selecting the  $k$  data points from the data, considering them as the centroids coordinate. Here  $k$  is the number of clusters.
- 2) Covariance Vector- The covariance vector is initially initialized by creating an identity matrix of size features\* features. The vector contains  $k$  number of identity matrices where  $k$  is the number of clusters.
- 3) Weight Vector- The weight vector is the weights of each cluster and initially all the clusters have same weight of  $1/k$  where  $k$  is number of clusters. So the vector contains  $k$  weights each with a value of  $1/k$ .

## Discussion of Maintaining $k$ Gaussians

Answer here...

The EM algorithm keeps on switching between an Expectation step and a Maximization step. In the Expectation step, the algorithm estimates the probability that each data point belongs to each component. In the Maximization step, the algorithm updates the parameters of the model to maximize the likelihood of the observed data given the estimated component probabilities. The GMM is a soft clustering method i.e the data point is not rigidly assigned to each cluster, instead, we have a posterior probability vector that stores the likelihood of each point going into a cluster. Hence  $k$  clusters are maintained at each iteration. For a data point, the maximum probability for going into a cluster as per the posterior probability calculation is selected. After every iteration, the weights and mean vectors for each cluster are updated.

## Discussion of Deciding Ties

Answer here... Here we maintain the posterior probability matrix for each data point and at every iteration. In case the probability of assigning a datapoint to cluster is same, I have used random selection, this means that in case of a tie between selecting a cluster, any conflicting cluster can be selected all having equal probability of getting selected.

## Discussion of Stopping Criteria

Answer here...

As per the algorithm provided, I have implemented a similar logic. When the mean vector is updated and the euclidean difference between the current centroid and the previous centroid is less than a particular threshold, in that case, the algorithm converges, and we get the final centroid coordinates, weight vector, probabilities, etc.

## Problem 2

Run your program,  $G_k$ , over the Diabetes data set and compare  $G_k$  with  $C_k$  (your  $k$ -means program from homework 4). Click on the below link to download the data set [50 points].

- [Diabetes 130-US Hospitals Data Set](#)

Answer the following questions:

1. Initialize  $G_k$  and  $C_k$  with the same set of initial points (initial centroids for  $C_k$  and  $\mu_i$ -s for  $G_k$  are identical) and run them for  $k = 2, \dots, 5$  for 20 runs each. Compare  $G_k$  and  $C_k$  using two different appropriate cluster validity techniques, i.e., internal, external or relative indices. Plots are generally a good way to convey complex ideas quickly, i.e., box plots, whisker plots. Discuss your results.

### R or Python script

```
# Sample R Script With Highlighting
```

```
# Sample Python Script With Highlighting
```

```
import pandas as pd
import swifter
5 import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
10 from scipy.stats import multivariate_normal
from sklearn.metrics import silhouette_score,
calinski_harabasz_score
from scipy.spatial.distance import euclidean

15 df= pd.read_csv('dataset_diabetes/diabetic_data.csv')

#Replacing Question Mark with NAN
import numpy as np
df.replace({'?':np.nan},inplace=True)
20 df.head()
```

```

encounter_id  patient_nbr  race gender  age  weight  admission_type_id  discharge_dispos
0      2278392    8222157  Caucasian Female  [0-10)   NaN    6      25      1      1      ...  No  No  N
1      149190     55629189  Caucasian Female  [10-20)   NaN    1      1      7      3      ...  No  Up  N
25 2      64410     86047875  AfricanAmerican  Female  [20-30)   NaN    1      1      1      7      2      ...  No  Up  No  N
3      500364     82442376  Caucasian Male  [30-40)   NaN    1      1      7      2      ...  No  Up  No  N
4      16680     42519267  Caucasian Male  [40-50)   NaN    1      1      7      1      ...  No  Steady  N
5 rows      50 columns
```

```

30 null_feature=[i for i in df.columns if
df[i].isnull().sum()>=1]
print('Null features {} \n'.format(null_feature))
print('Feature \t null_count \t not_null_count')
35 for i in null_feature:
    print('{} \t {} \t\t
    {}'.format(i,df[i].isnull().sum(),df[i].count()))

Null features ['race', 'weight', 'payer_code', 'medical_specialty', 'diag_1', 'diag_2', 'diag_3']

40 Feature      null_count      not_null_count
race          2273          99493
weight        98569          3197
payer_code    40256          61510
45 medical_specialty  49949          51817
diag_1         21          101745
diag_2         358          101408
diag_3         1423          100343

50 #Dropping columns with count of null values around the
count of not null values. As they dont provide
significant information or
mostly contains null data
55 df.drop(['weight','payer_code','medical_specialty'],axis=
1,inplace=True)

60 df['age'] = df['age'].replace({'[0-10]': 5, '[10-20]':
15, '[20-30]': 25, '[30-40]': 35, '[40-50]': 45, '[50-
60]': 55, '[60-70]': 65, '[70-80]': 75, '[80-90]': 85,
'[90-100]': 95})

65 df['readmitted'] = df['readmitted'].replace({'>30':1,'<30':1,'NO':0})

for i in df.select_dtypes(include=['int',
'float']).columns.to_list():
    print('The numeric feature is {} \n The value counts
70 are {}'.format(i,df[i].value_counts()) )

df.drop(columns=['encounter_id','patient_nbr'],inplace=True,axis=1)

75 meds=['max_glu_serum', 'AlCresult', 'metformin',
'repaglinide', 'nateglinide', 'chlorpropamide',
'glimepiride', 'acetoexamide', 'glipizide',
'glyburide', 'tolbutamide', 'pioglitazone',
'rosiglitazone',
'acarbose', 'miglitol', 'troglitazone',
80 'tolazamide', 'examide', 'citoglipton', 'insulin',
'glyburide-metformin',
'glipizide-metformin', 'glimepiride-pioglitazone',

```

```

    'metformin-rosiglitazone', 'metformin-pioglitazone']
df_value_counts=pd.DataFrame()
85 for i in meds:
    value_counts=df[i].value_counts()
    percent = []
    for j in value_counts.index:
        percent.append(value_counts[j] *100/ len(df))
90    ## PErcentage dataframe to store the feature, its
    unique values, the count and the percentage
    df_temp=pd.DataFrame({'Feature':i,'Value':
        value_counts.index, 'Count': value_counts.values,
        'Percentage': percent})
95    df_value_counts=pd.concat ([df_value_counts,df_temp],I
        gnore_index=True)
df_value_counts.head(80)

Feature  Value      Count      Percentage
100 0    max_glu_serum  None 96420      94.746772
1    max_glu_serum  Norm 2597  2.551933
2    max_glu_serum  >200 1485  1.459230
3    max_glu_serum  >300 1264  1.242065
4    AlCresult  None 84748      83.277322
105 ...    ...    ...    ...
74    glimepiride-pioglitazone  Steady  1    0.000983
75    metformin-rosiglitazone  No  101764  99.998035
76    metformin-rosiglitazone  Steady  2    0.001965
77    metformin-pioglitazone  No  101765  99.999017
110 78    metformin-pioglitazone  Steady  1    0.000983
79 rows    4 columns

skewed_data=df_value_counts[df_value_counts['Percentage']
>95]['Feature'].to_list()
115 df.drop(columns=skewed_data,inplace=True)

label_encoded_columns=[]
for i in df.select_dtypes(include=
['object']).columns.to_list():
120     if i not in skewed_data:
        label_encoded_columns.append(i)
## The columns remaining after all the EDA that are to be label encoded

125 df_cleaned_dia=df.copy()
## creating the copy of data before performing label encoding and one hot encoding
one_hot = pd.get_dummies(df_cleaned_dia[['gender','race']])
label_encoded_columns.remove('diag_1')
label_encoded_columns.remove('diag_2')
130 label_encoded_columns.remove('diag_3')
label_encoded_columns.remove('gender')
label_encoded_columns.remove('race')
# combine the one-hot encoded columns with the original dataframe
df_cleaned_dia = pd.concat([df_cleaned_dia, one_hot], axis=1)
135 df_cleaned_dia.drop(columns=['diag_1','diag_2','diag_3','gender','race'],inplace=True)

```

```

##
df_cleaned_dia[label_encoded_columns]=df_cleaned_dia[label_encoded_columns].swifter.apply(LabelEncoder().fit_transform)
140

from sklearn.metrics import silhouette_score,
145 calinski_harabasz_score
def GMM_initialization(df,k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
    means_matrix = df.sample(n=k).values
    150 identity_matrix=np.eye(number_of_columns)
    covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)
    return
    155 means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns

def
calculate_posterior(data,means_matrix,covariance_matrix,weights_matrix,k,number_of_columns):
160
    posterior=np.zeros(k)
    for i in range(k):
        try:
            pseudo_inverse =
            165 np.linalg.pinv(covariance_matrix[i] +
            np.eye(covariance_matrix[i].shape[0]) * 1e-2,
            rcond=1e-10)
            posterior[i] = multivariate_normal.pdf(data,
            170 mean=means_matrix[i], cov=pseudo_inverse)
        except Exception as e:
            continue
    return posterior*weights_matrix/(posterior*weights_matrix).sum()

175 def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
    return_counts=True)
    missing_labels=[i for i in range(k) if i not in
    unique_values]
    180 unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

    indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
    185 size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
        labels[val]=missing_labels[i]
    return labels

```

```

190 def sum_of_square_error_em(new_centroids, data, labels):
    columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
195 # Rename the '0' column of the labels dataframe to
    'Label'
    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
200 its assigned centroid
    for i in range(len(new_centroids)):
        distance =
            np.sum(np.square(data[data['Label']==i]
                [columns] - new_centroids.iloc[i]
205 [columns],dtype=np.float64), axis=1)
            #print(distance)
            sse.append(distance.sum())
    # Return the sum of squared errors

210 a=sum(sse)
    return a

def Calinski_index_em(df_data, clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
215 return ch_score

def GMM(df_cleaned_dia,k,tao):
220 scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.colu
225 mns)

    means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns=
    GMM_initialization(scaled_input_df,k)
    likelihood=0
230 means_matrix_initial=means_matrix
    for i in range(k):
        try:

            pseudo_inverse =
235 np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0])
                )* 1e-10))
            likelihood=likelihood+weights_matrix[i]*multi
            variate_normal.logpdf(scaled_input,means_matrix
                [i], pseudo_inverse)
240 except Exception as e:
            continue

```

```

log_likelihood_old=np.sum(likelihood)
old_means_matrix_df=pd.DataFrame(means_matrix)
posterior_probability = np.zeros((scaled_input.shape[0], k))
245 iterations=0

while (True):
    iterations+=1
    # Expectation
250 for i in range(scaled_input.shape[0]):
        posterior_probability[i] =
            calculate_posterior(scaled_input[i],
                                means_matrix, covariance_matrix, weights_matrix,
                                k, number_of_columns)

255 # Maximization
posterior_probability=np.nan_to_num(posterior_probability, nan=0)
for i in range(k):
260     # Calculating weight
    weight = posterior_probability[:, i].sum()
    #print(weight)
    # Updating each centroid
    means_matrix[i] = (posterior_probability[:,
265 i] @ scaled_input) / weight
    #print(1,means_matrix[i])
    # Subtracting the mean value from data
    scaled_input_diff = scaled_input - means_matrix[i]

270 # Update the covariance matrix
    covariance_matrix[i] =
        (posterior_probability[:, i] *
         scaled_input_diff.T @ scaled_input_diff) /
        weight

275 # Update the weights matrix
    weights_matrix[i] = weight / number_of_rows

280 likelihood=0
for i in range(k):
    try:
        pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
285 np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
        likelihood=likelihood+weights_matrix[i]*m
        ultivariate_normal.logpdf(scaled_input, means_matrix[i], pseudo_inverse)
    except Exception as e:
290         continue
log_likelihood_new =np.sum(likelihood)

new_means_matrix_df=pd.DataFrame(means_matrix)
distance = []

```

```

295     for col in new_means_matrix_df.columns:
        col_distance =
            euclidean(old_means_matrix_df[col],
                new_means_matrix_df[col])
            distance.append(col_distance)
300     tao_calculated=sum(distance)/k

    if tao_calculated<
305     tao:#log_likelihood_new>log_likelihood_old and
        100*((log_likelihood_new - log_likelihood_old) /
            log_likelihood_old)<tao:

        print("Converged")
310         labels=np.argmax(posterior_probability,axis=1
            )
            labels=maintain_k_clusters(labels,k)
            labels_df=pd.DataFrame(labels)
            means_matrix_df=pd.DataFrame(means_matrix,col
315             umns=scaled_input_df.columns)
            sse=sum_of_square_error_em(means_matrix_df,
                scaled_input_df, labels_df)
            clainski=
                Calinski_index_em(scaled_input_df,labels_df)
320         return sse,clainski,means_matrix_initial
    #else:
        #log_likelihood_old=log_likelihood_new

    if iterations>100:
325         print("Max iteration reached")
            labels=np.argmax(posterior_probability,axis=1
                )
                labels=maintain_k_clusters(labels,k)
                labels_df=pd.DataFrame(labels)
330                means_matrix_df=pd.DataFrame(means_matrix,col
                    umns=scaled_input_df.columns)
                    sse=sum_of_square_error_em(means_matrix_df,
                        scaled_input_df, labels_df)
                    clainski=
335                        Calinski_index_em(scaled_input_df,labels_df)
                        return sse,clainski,means_matrix_initial

import time
340 from scipy.spatial.distance import euclidean
def initialize_centroids(df, k,means_matrix):
    """
    Function to initialize random centroids from dataset.
    Input:
345     - df: pandas dataframe with the data
        - k: integer number of clusters
    Output:

```



```

    - temp_df: pandas dataframe with the centroids as columns and index as label
    """
350
    centroids=pd.DataFrame(means_matrix,columns=df.columns)
    centroids=centroids.T
355    centroids.index.name = 'Label'
    return centroids

360 def assign_labels(df, centroids):
    """
    Function to calculate the closest centroid label for each row in a dataframe.
    Input:
        - df: pandas dataframe with the data
365        - centroids: pandas dataframe with the centroids as columns and index as label
    Output:
        - distances.idxmin(axis=1): pandas series with
          the label of the closest centroid for each row
          in df
370    """
    distances = centroids.swifter.apply(lambda x:
    np.sqrt(((df - x) ** 2).sum(axis=1))) # Calculate
    the Euclidean distance between each row in df and
    each centroid
375    return distances.idxmin(axis=1) # Get the index of
    the minimum distance, which corresponds to the label
    of the closest centroid

380 def new_centroids(df_label, df1):
    """
    Function to calculate the new centroids based on the
    current labels of the rows.
    Input:
385        - df_label: pandas series with the label of the
          closest centroid for each row in df1
        - df1: pandas dataframe with the data
    Output:
        - new_centroids.T: pandas dataframe with the new
390        centroids as columns and index as feature name
    """
    joined_df = df1.join(df_label)
    joined_df.rename(columns={0: 'Label'}, inplace=True) # Rename the column with the label
    # Calculate the mean of the rows with the same label
395    return joined_df.groupby('Label').mean().T #
    Transpose the dataframe to have the new centroids as
    columns and index as feature name

400

```

```

def error_clusters(df_new_centroids, df1, df_label):
    """
    Calculate the error rate of each cluster.

    Args:
    - df_label (pandas.DataFrame): the label of the
    nearest centroid for each data point.
    - df1 (pandas.DataFrame): the dataset.
    - df_new_centroids (pandas.DataFrame): The new centroids computed in the current iteration.

    Returns:
    - error_rate (float): the total error rate of all clusters.
    """

    # Calculate mean value
    mean_centroid = df1.groupby('readmitted').mean().reset_index()
    # Transpose the new centroids dataframe and reset the index
    new_centroids = df_new_centroids.T
    # Get the columns of the data dataframe
    columns = df1.columns

    sse = []
    # Compute the distance between each data point and its assigned centroid
    for i in range(len(new_centroids)):  ### centroid
        s = []
        for j in range(len(mean_centroid)):  ### mean centroid
            # Compute the distance between each data point and its assigned centroid
            distance =
            np.sum(np.square(mean_centroid[mean_centroid[
                'readmitted'] == j][columns] -
                new_centroids.iloc[i][columns]), axis=1)
            s.append(distance.iloc[0])
        sse.append(s)
    ## key is the cluster number and value is the merged value
    merge_label = pd.DataFrame(sse).idxmin(axis=1).to_dict()
    ## Merging cluster based on the target variable
    df_label[0] = df_label[0].replace(merge_label)

    df1 = df1.join(df_label) # add the label column to the dataset
    df1.rename(columns={0: 'Label'}, inplace=True) # rename the label column
    error_list = []
    for i in df1['Label'].value_counts().index:
        df_cluster = df1[df1['Label'] == i] # filter the
        dataset to include only the data points in the
        current cluster
        y = len(df_cluster[df_cluster['readmitted'] ==
            1]) # count the number of data points in the
        current cluster that were readmitted
        n = len(df_cluster[df_cluster['readmitted'] ==
            0]) # count the number of data points in the
        current cluster that were not readmitted
        if y == 0 and n == 0:

```

```

    error = 0
455     else:
        error = n / (n + y) # calculate the error
        rate of the current cluster
        error_list.append(error)
    return round(sum(error_list),4)
460

def sum_of_square_error(new_centroids, data, labels):
    """
    Computes the sum of squared errors between the data
465     points and their assigned centroids.

    Args:
    new_centroids (DataFrame): The new centroids computed in the current iteration.
    data (DataFrame): The input data points.
470     labels (DataFrame): The labels assigned to each data point.

    Returns:
    The sum of squared errors.
    """
    # Transpose the new centroids dataframe and reset the index
    new_centroids = new_centroids.T.reset_index()
    # Get the columns of the data dataframe
    columns = new_centroids.columns
    # Join the data dataframe and the labels dataframe
480     data = data.join(labels)
    # Rename the '0' column of the labels dataframe to 'Label'
    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
485     its assigned centroid
    for i in range(len(new_centroids)):
        distance =
            np.sum(np.square(data[data['Label']==i][columns]
                - new_centroids.iloc[i][columns]), axis=1)
490         sse.append(sum(distance))
    # Return the sum of squared errors
    return sum(sse)

def Calinski_index(df_data, clusters):
495     ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score

500

def kmeans_lyod_with_error(df1, k, tou, means_matrix_initial):
    """
    Function to run the K-means Lloyd algorithm.
    Input:
505     - df1: pandas dataframe with the data
        - k: integer number of clusters

```

```

    - tou: float tolerance level to stop the algorithm
Output:
    - centroids: pandas dataframe with the final centroids as columns and index as label
510 """
start_time=time.time()
centroids = initialize_centroids(df1,
k,means_matrix_initial) # Initialize random centroids
initial_list_of_columns = centroids.columns.to_list()
515 iteration = 0
while True:
    # Assign labels to current centroids
    df_label = assign_labels(df1, centroids)
    df_label = pd.DataFrame(df_label)
520 # Calculate new centroids
    df_new_centroids = new_centroids(df_label, df1)
    new_list_of_columns =
    df_new_centroids.columns.to_list()
    # Keep the number of clusters the same i.e
525 maintain same k
    for i in initial_list_of_columns:
        if i not in new_list_of_columns:
            df_new_centroids[i] = centroids[i]
    # Calculate tao
530 distance = []
    for col in centroids.columns:
        col_distance = euclidean(centroids[col], df_new_centroids[col])
        distance.append(col_distance)
    tao_calculated=sum(distance)/k #Used the formula provided for calculating Tao
535 sse = sum_of_square_error(df_new_centroids, df1, df_label)
    #error=error_clusters(df_label,df1,k)
    end_time= time.time()
    clainski= Calinski_index(df1,df_label)
    if iteration>100:
540     print("Iteration exceeded")

    return sse,clainski
    break

545 if tao_calculated<tou or iteration >100: #if
the convergence is met, kmeans will stop or
else if the convergence is never met, after 100
iteration code will stop
    return sse,clainski
550     break # otherwise indefinite loop
else:
    centroids= df_new_centroids # In case we
    need more iterations, the centroids
    calculated at this step acts as input
555 iteration+=1

scaler = StandardScaler()

```

```

560 scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

565 error_matrix_em=[]
    error_matrix_kmeans=[]
    for i in range(2,6):
        for j in range(1,21):
570             sse,clainski,means_matrix_initial=GMM(df_cleaned_
                dia,i,10)
                error_matrix_em.append([i,sse,clainski])

                sse,clainski=kmeans_lyod_with_error(scaled_input_
575                 df,i,10,means_matrix_initial)
                error_matrix_kmeans.append([i,sse,clainski])
    error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse','clainski'])
    error_df_kmeans= pd.DataFrame(error_matrix_kmeans,columns=
        ['number_of_cluster', 'sse','clainski'])

580

    error_df_em.to_csv('6_em.csv',index=False)
    error_df_kmeans.to_csv('6_kmeans.csv',index=False)

585

    error_df_em['algo']='em'
    error_df_kmeans['algo']='kmeans'

590

    run_time_diab=pd.DataFrame()
    run_time_diab=pd.concat( [
    error_df_em[['algo','number_of_cluster','sse','clainski']],
595     error_df_kmeans[['algo','number_of_cluster',
        'sse','clainski']]
        ],ignore_index=True )

    import seaborn as sns

600

    fig, ax = plt.subplots(figsize=(8,6))

    sns.boxplot(x='number_of_cluster', y='sse', hue='algo',

605             data=run_time_diab[run_time_diab['algo'].isin
                (['kmeans','em'])],ax=ax);
    plt.title('Box Plot of SSE for GMM and K means')
    plt.show()

610

    import seaborn as sns

```

```


fig, ax = plt.subplots(figsize=(8,6))

615 sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
               data=run_time_diab[run_time_diab['algo'].isin(['kmeans','em'])],ax=ax);
plt.title('Box Plot of Clainski Index for GMM and K means')
plt.show()

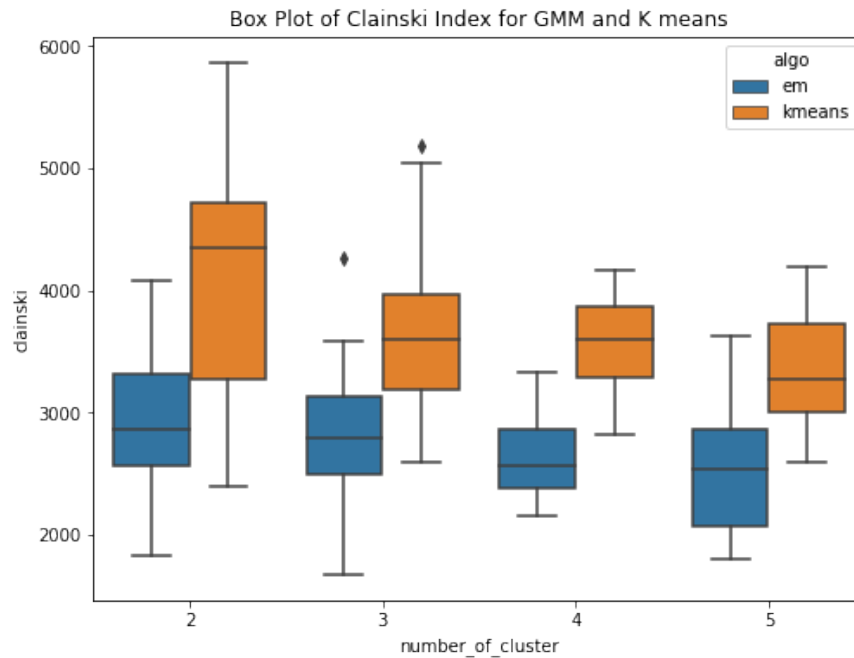
```

## Plot/s

Place images here with suitable captions.



(3)



(4)

## Discussion of Experiments

Answer here. . .

The GMM (EM) algorithm and K-means algorithm is ran over the diabetes dataset 20 times each for clusters ranging from 2 to 5. Have used within sum of square error and Calinski-Harabasz score to evaluate the clustering techniques. Within Sum of Squares (WSS), which calculates the total sum of distances between each point in a cluster and its centroid. Calinski-Harabasz score (CHS), which evaluates the quality of the clustering solution by comparing the separation between clusters to the dispersion within clusters. From the box plots, it can be seen that the median within the sum of square error for k means is less than the median within sse of EM. This means for the given dataset, K means have performed better or the clusters are well separated and the within-cluster the data points are very close to the centroid. Whereas for GMM, it has higher sse, one of the reasons could be the clusters formed are overlapping and this may lead to higher error. This can be confirmed by the Calinski Harabasz score (CHS). The median CHS score for kmeans is higher than that of the GMM for all the cluster, showing the clusters formed by Kmeans are well seperated. We can use other techniques to perform the validity and this highly depends upon the requirements and use case.

2. Run your  $G_k$  without updating the covariance matrices and priors across iterations. Compare  $G_k$  and  $C_k$  using two different appropriate cluster validity techniques, i.e., internal, external or relative indices. Plots are generally a good way to convey complex ideas quickly, i.e., box plots, whisker plots. Discuss your results.

## R or Python script

```
# Sample R Script With Highlighting
```

```
import pandas as pd
import swifter
import matplotlib.pyplot as plt
5 import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
from scipy.stats import multivariate_normal
10 from sklearn.metrics import silhouette_score, calinski_harabasz_score
from scipy.spatial.distance import euclidean

def GMM_initialization(df,k):
15     number_of_rows=df.shape[0]
     number_of_columns=df.shape[1]
     means_matrix = df.sample(n=k).values
     identity_matrix=np.eye(number_of_columns)
     covariance_matrix=np.array([identity_matrix]*k)
20     weights_matrix=np.array([float(1/k)]*k)
     return means_matrix,covariance_matrix,weights_matrix,number_
           of_rows,number_of_columns

25 def calculate_posterior(data,means_matrix,covariance_matrix,w
           eights_matrix,k,number_of_columns):
     posterior=np.zeros(k)
     for i in range(k):
         try:
30             pseudo_inverse =
                 np.linalg.pinv(covariance_matrix[i] +
                 np.eye(covariance_matrix[i].shape[0]) * 1e-
                 2, rcond=1e-10)
                 posterior[i] = multivariate_normal.pdf(data,
35                 mean=means_matrix[i], cov=pseudo_inverse)
         except Exception as e:
             continue
     return
     posterior*weights_matrix#/(posterior*weights_matrix).
40     sum()

def maintain_k_clusters(labels,k):
     unique_values, value_counts = np.unique(labels,
```



```

    return_counts=True)
45 missing_labels=[i for i in range(k) if i not in
    unique_values]
    unique_values_to_be_replaced=[unique_values[i] for i
    in np.where(value_counts > 1)[0]]

50 indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
    size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
55     labels[val]=missing_labels[i]
    return labels

def sum_of_square_error_em(new_centroids, data, labels):
60     columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
    # Rename the '0' column of the labels dataframe to
    'Label'
65     data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
    its assigned centroid
    for i in range(len(new_centroids)):
70         distance =
            np.sum(np.square(data[data['Label']==i][columns]
            - new_centroids.iloc[i]
            [columns],dtype=np.float64), axis=1)
            #print(distance)
75         sse.append(distance.sum())
    # Return the sum of squared errors

    a=sum(sse)
    return a
80

def Calinski_index_em(df_data,clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score
85

def GMM(df_cleaned_dia,k,tao):
    scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
90     scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.colu
    mns)

95     means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns=
    GMM_initialization(scaled_input_df,k)

```

```

likelihood=0
means_matrix_initial=means_matrix
for i in range(k):
    100     try:

        pseudo_inverse =
        np.linalg.pinv(covariance_matrix[i] +
        np.diag(np.ones(covariance_matrix[i].shape[0]
    105     ) * 1e-10))
        likelihood=likelihood+weights_matrix[i]*multi
        variate_normal.logpdf(scaled_input,means_matr
        ix[i], pseudo_inverse)
    except Exception as e:
    110     continue

log_likelihood_old=np.sum(likelihood)
old_means_matrix_df=pd.DataFrame(means_matrix)
posterior_probability =
np.zeros((scaled_input.shape[0], k))
    115     iterations=0

    while (True):
        iterations+=1
        # Expectation
    120     for i in range(scaled_input.shape[0]):
        posterior_probability[i] =
        calculate_posterior(scaled_input[i],
        means_matrix,covariance_matrix,weights_matrix
        ,k,number_of_columns)

    125     # Maximization
        posterior_probability=np.nan_to_num(posterior_pro
        bability, nan=0)
        for i in range(k):
    130             # Calculating weight
            weight = posterior_probability[:, i].sum()
            #print(weight)
            # Updating each centroid
            means_matrix[i] = (posterior_probability[:, i] @ scaled_input) / weight
    135             #print(1,means_matrix[i])
            # Subtracting the mean value from data
            scaled_input_diff = scaled_input - means_matrix[i]

            # Update the covariance matrix
    140             #covariance_matrix[i] = (posterior_probability[:, i] *
            scaled_input_diff.T @ scaled_input_diff) / weight

            # Update the weights matrix
            #weights_matrix[i] = weight / number_of_rows

    145

        likelihood=0
        for i in range(k):
            try:

```

```

150         pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
            np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*m
155         multivariate_normal.logpdf(scaled_input, means_matrix[i], pseudo_inverse)
            except Exception as e:
                continue
        log_likelihood_new = np.sum(likelihood)

160
        new_means_matrix_df=pd.DataFrame(means_matrix)
        distance = []
        for col in new_means_matrix_df.columns:
            col_distance =
165             euclidean(old_means_matrix_df[col],
                new_means_matrix_df[col])

            distance.append(col_distance)
        tao_calculated=sum(distance)/k

170

        if tao_calculated<
        tao:#log_likelihood_new>log_likelihood_old and
175         100*((log_likelihood_new - log_likelihood_old) /
            log_likelihood_old)<tao:

            print("Converged")
            labels=np.argmax(posterior_probability,axis=1)
180         )
            labels=maintain_k_clusters(labels,k)
            labels_df=pd.DataFrame(labels)
            means_matrix_df=pd.DataFrame(means_matrix, columns=scaled_input_df.columns)
            sse=sum_of_square_error_em(means_matrix_df,
185             scaled_input_df, labels_df)
            clainski=
            Calinski_index_em(scaled_input_df, labels_df)
            return sse,clainski,means_matrix_initial
190     #else:
        #log_likelihood_old=log_likelihood_new

        if iterations>100:
            print("Max iteration reached")
195             labels=np.argmax(posterior_probability,axis=1)
            )
            labels=maintain_k_clusters(labels,k)
            labels_df=pd.DataFrame(labels)
            means_matrix_df=pd.DataFrame(means_matrix, columns=scaled_input_df.columns)
200             sse=sum_of_square_error_em(means_matrix_df,
                scaled_input_df, labels_df)

```

```

        clainski=
        Calinski_index_em(scaled_input_df, labels_df)
205         return sse, clainski, means_matrix_initial

import time
210 from scipy.spatial.distance import euclidean
def initialize_centroids(df, k, means_matrix):
    """
    Function to initialize random centroids from dataset.
    Input:
215     - df: pandas dataframe with the data
        - k: integer number of clusters
    Output:
        - temp_df: pandas dataframe with the centroids as columns and index as label
    """

    centroids=pd.DataFrame(means_matrix, columns=df.columns)
    centroids=centroids.T
225     centroids.index.name = 'Label'
    return centroids

230 def assign_labels(df, centroids):
    """
    Function to calculate the closest centroid label for each row in a dataframe.
    Input:
235     - df: pandas dataframe with the data
        - centroids: pandas dataframe with the centroids as columns and index as label
    Output:
        - distances.idxmin(axis=1): pandas series with
          the label of the closest centroid for each row
          in df
240     """
    distances = centroids.swifter.apply(lambda x:
    np.sqrt(((df - x) ** 2).sum(axis=1))) # Calculate
    the Euclidean distance between each row in df and
    each centroid
245     return distances.idxmin(axis=1) # Get the index of
    the minimum distance, which corresponds to the label
    of the closest centroid

250 def new_centroids(df_label, df1):
    """
    Function to calculate the new centroids based on the
    current labels of the rows.
    Input:
255     - df_label: pandas series with the label of the

```

```

        closest centroid for each row in df1
        - df1: pandas dataframe with the data
    Output:
        - new_centroids.T: pandas dataframe with the new
260         centroids as columns and index as feature name
        """
        joined_df = df1.join(df_label)
        joined_df.rename(columns={0: 'Label'}, inplace=True) # Rename the column with the label
        # Calculate the mean of the rows with the same label
265         return joined_df.groupby('Label').mean().T #
        Transpose the dataframe to have the new centroids as
        columns and index as feature name

270 def error_clusters(df_new_centroids, df1, df_label):
        """
        Calculate the error rate of each cluster.

275         Args:
        - df_label (pandas.DataFrame): the label of the
        nearest centroid for each data point.
        - df1 (pandas.DataFrame): the dataset.
        - df_new_centroids (pandas.DataFrame): The new centroids computed in the current iteration.

280         Returns:
        - error_rate (float): the total error rate of all clusters.
        """

285         #Calculate mean value
        mean_centroid=df1.groupby('readmitted').mean().reset_index()
        # Transpose the new centroids dataframe and reset the index
        new_centroids= df_new_centroids.T
290         # Get the columns of the data dataframe
        columns = df1.columns

        sse = []
        # Compute the distance between each data point and its assigned centroid
295         for i in range(len(new_centroids)): #### centroid
            s=[]
            for j in range(len(mean_centroid)): ### mean centroid
                # Compute the distance between each data point and its assigned centroid
                distance =
300                 np.sum(np.square(mean_centroid[mean_centroid[
                    'readmitted']==j][columns] -
                    new_centroids.iloc[i][columns]), axis=1)
                s.append(distance.iloc[0])
            sse.append(s)
305         ## key is the cluster number and value is the merged value
        merge_label=pd.DataFrame(sse).idxmin(axis=1).to_dict()
        ## Merging cluster based on the target variable
        df_label[0]=df_label[0].replace(merge_label)

```

```

310 df1 = df1.join(df_label) # add the label column to the dataset
df1.rename(columns={0: 'Label'}, inplace=True) # rename the label column
error_list = []
for i in df1['Label'].value_counts().index:
    df_cluster = df1[df1['Label'] == i] # filter the
315 dataset to include only the data points in the
    current cluster
    y = len(df_cluster[df_cluster['readmitted'] ==
1]) # count the number of data points in the
    current cluster that were readmitted
320 n = len(df_cluster[df_cluster['readmitted'] ==
0]) # count the number of data points in the
    current cluster that were not readmitted
    if y == 0 and n == 0:
        error = 0
325 else:
        error = n / (n + y) # calculate the error
        rate of the current cluster
        error_list.append(error)
return round(sum(error_list),4)
330

def sum_of_square_error(new_centroids, data, labels):
    """
    Computes the sum of squared errors between the data
335 points and their assigned centroids.

    Args:
    new_centroids (DataFrame): The new centroids computed in the current iteration.
    data (DataFrame): The input data points.
340 labels (DataFrame): The labels assigned to each data point.

    Returns:
    The sum of squared errors.
    """
    # Transpose the new centroids dataframe and reset the index
    new_centroids = new_centroids.T.reset_index()
    # Get the columns of the data dataframe
    columns = new_centroids.columns
    # Join the data dataframe and the labels dataframe
350 data = data.join(labels)
    # Rename the '0' column of the labels dataframe to 'Label'
    data.rename(columns={0: 'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
    its assigned centroid
355 for i in range(len(new_centroids)):
        distance =
        np.sum(np.square(data[data['Label']==i][columns]
        - new_centroids.iloc[i][columns]), axis=1)
360 sse.append(sum(distance))
    # Return the sum of squared errors

```

```

    return sum(sse)

def Calinski_index(df_data, clusters):
365     ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score

370
def kmeans_lyod_with_error(df1, k, tou, means_matrix_initial):
    """
    Function to run the K-means Lloyd algorithm.
    Input:
375     - df1: pandas dataframe with the data
     - k: integer number of clusters
     - tou: float tolerance level to stop the algorithm
    Output:
     - centroids: pandas dataframe with the final centroids as columns and index as label
380     """
    start_time=time.time()
    centroids = initialize_centroids(df1,
    k, means_matrix_initial) # Initialize random centroids
    initial_list_of_columns = centroids.columns.to_list()
385    iteration = 0
    while True:
        # Assign labels to current centroids
        df_label = assign_labels(df1, centroids)
        df_label = pd.DataFrame(df_label)
390        # Calculate new centroids
        df_new_centroids = new_centroids(df_label, df1)
        new_list_of_columns =
        df_new_centroids.columns.to_list()
        # Keep the number of clusters the same i.e
395        maintain same k
        for i in initial_list_of_columns:
            if i not in new_list_of_columns:
                df_new_centroids[i] = centroids[i]
        # Calculate tao
400        distance = []
        for col in centroids.columns:
            col_distance = euclidean(centroids[col], df_new_centroids[col])
            distance.append(col_distance)
        tao_calculated=sum(distance)/k #Used the formula provided for calculating Tao
405        sse = sum_of_square_error(df_new_centroids, df1, df_label)
        #error=error_clusters(df_label,df1,k)
        end_time= time.time()
        clainski= Calinski_index(df1,df_label)
        if iteration>100:
410            print("Iteration exceeded")

            return sse, clainski
            break

```

```

415     if tao_calculated<tou or iteration >100:    #if
the convergence is met, kmeans will stop or
else if the convergence is never met, after 100
iteration code will stop
        return sse,clainski
420     break                                     # otherwise indefinite loop
    else:
        centroids= df_new_centroids # In case we
        need more iterations, the centroids
        calculated at this step acts as input
425     iteration+=1

scaler = StandardScaler()
430 scaler.fit(df_cleaned_dia)
scaled_input=scaler.transform(df_cleaned_dia)

scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

435 error_matrix_em=[]
error_matrix_kmeans=[]
for i in range(2,6):
    for j in range(1,21):
440         sse,clainski,means_matrix_initial=GMM(df_cleaned_
dia,i,10)
        error_matrix_em.append([i,sse,clainski])

        sse,clainski=kmeans_lyod_with_error(scaled_input_
445         df,i,10,means_matrix_initial)
        error_matrix_kmeans.append([i,sse,clainski])
error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse','clainski'])
error_df_kmeans= pd.DataFrame(error_matrix_kmeans,columns=
['number_of_cluster', 'sse','clainski'])
450

error_df_em.to_csv('6_em_no_update.csv',index=False)
error_df_em.to_csv('6_kmeans_no_update.csv',index=False)

455 error_df_em['algo']='em'
error_df_kmeans['algo']='kmeans'

run_time_diab=pd.DataFrame()
run_time_diab=pd.concat( [ error_df_em[['algo','number_of_cluster','sse','clainski']],
error_df_kmeans[['algo','number_of_cluster', 'sse','clainski']]
],ignore_index=True )

465 import seaborn as sns

fig, ax = plt.subplots(figsize=(8,6))

```



```

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
470         data=run_time_diab[run_time_diab['algo'].isin(['kmeans','em']),ax=ax);
plt.title('Box Plot of SSE for GMM and K means without updating the covariance and priors')
plt.show()

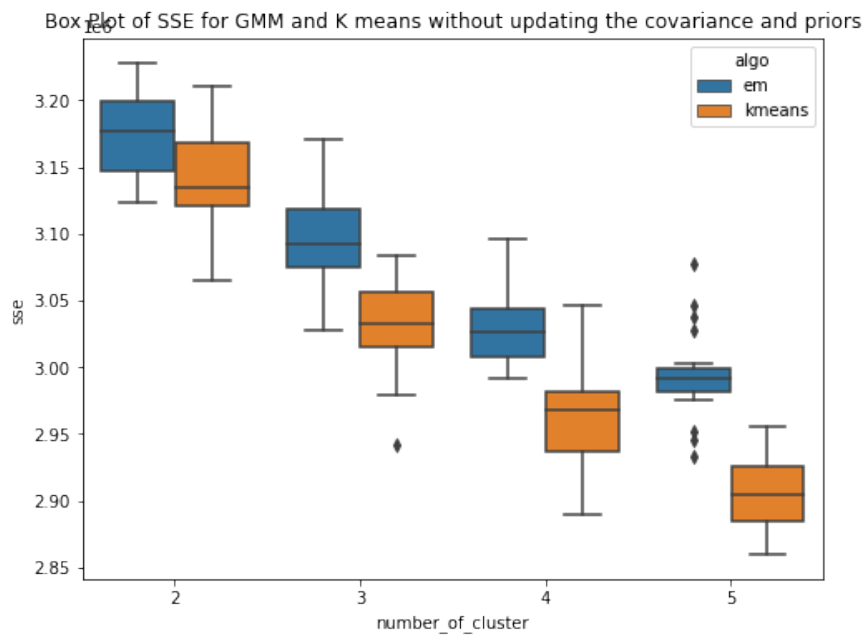
import seaborn as sns
475
fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin(['kmeans','em']),ax=ax];
480 plt.title('Box Plot of Clainski for GMM and K means without updating the covariance and priors')
plt.show()

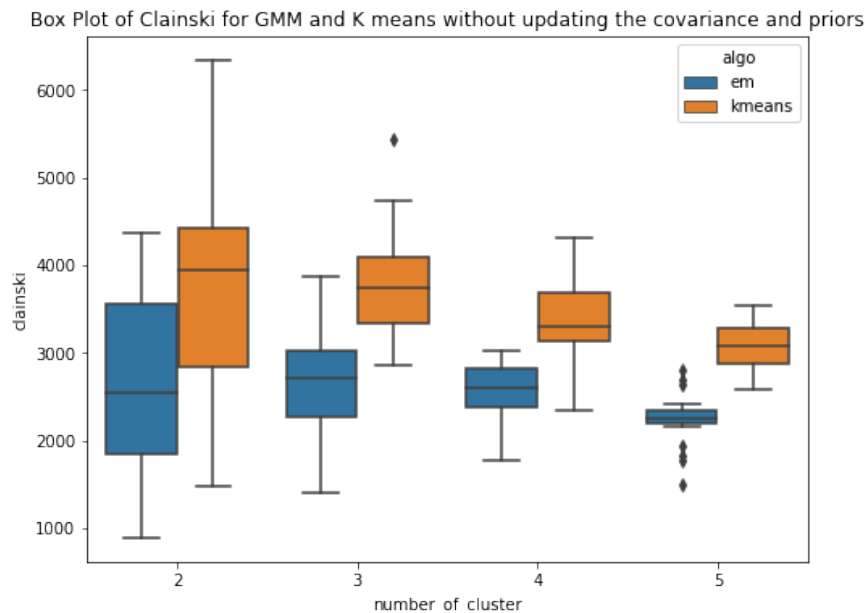
```

## Plot/s

Place images here with suitable captions.



(5)



(6)

## Discussion of Experiments

Answer here. . .

If the GMM algorithm doesn't update the covariances and weights vector, K-means and GMM can converge at the same time. This is because when the covariance matrix is kept fixed, the GMM algorithm becomes similar to K-means, which assigns data points to the nearest centroid based on the Euclidean distance. When the covariances are fixed in GMM, the algorithm still estimates the mean values and mixing coefficients (weights vector) using the EM algorithm. In the E-step, the algorithm calculates the probability of each data point belonging to each Gaussian distribution based on their mean values and fixed covariance matrix. The M-step then updates the mean values and mixing coefficients based on the calculated probabilities. However, this process resembles the K-means algorithm, and both algorithms may converge at nearly similar points if the covariances and weights vector are not updated during the GMM algorithm. This can be seen from the box plot of within sum of square error. The median within SSE of kmeans is slightly less than that of the GMM, but it is comparable. The errors are almost similar, showing convergence. Also the box plot of CHS score shows kmeans and GMM have almost similar values with kmeans more on the positive side of the graph, conveying the same story of convergence when the weights and covariance vectors are not updated.

3. Perform PCA over the Diabetes data set. Create a new data set,  $\Delta_R$ , with using 90% of the variance. Compare  $G_k$  and  $C_k$  over  $\Delta_R$  using two different appropriate cluster validity techniques, i.e., internal, external or relative indices. Plots are generally a good way to convey complex ideas quickly, i.e., box plots, whisker plots. Discuss your results.

## R or Python script

```
# Sample R Script With Highlighting
```

```

# Sample Python Script With Highlighting

import pandas as pd
import swifter
5 import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
import numpy as np
10 from sklearn.preprocessing import StandardScaler
from numpy import linalg as LA
import warnings

15
def PCA(df, threshold):
    ## Performing Standardization so that all features
    are given same importance initially and each feature
    can contribute
    20 ## equally to PC irrespective of scale or magnitude
    df = pd.DataFrame(StandardScaler().fit_transform(df))
    ## Performing Data Centering on standardized
    dataframe
    centred_df = df - np.mean(df, axis=0)
    25 ## Calculating Covariance
    covariance = np.cov(df.T)
    eigen_values, eigen_vectors = LA.eig(covariance)
    ## Sorting the eigen values in descending order,
    using argsort, we get the indices of eigen values in
    30 descending order
    sorted_index = eigen_values.argsort()[::-1]
    df_variance = pd.DataFrame(eigen_values[sorted_index] / sum(eigen_values), columns=['variance'])
    df_variance['cumulative_variance'] =
    35 df_variance['variance'].cumsum()
    ## Number of principal components required to cover
    variance uptill certian threshold
    df_number_pc_var = df_variance[df_variance['cumulative_
    variance'] <= threshold]
    40 number_of_pc = len(df_number_pc_var)
    print("The number of principal components required
    to cover {} percent variance are
    {}".format(threshold*100, number_of_pc))
    ## Selecting the required number of eigen vectores
    45 for performing dot product
    selected_eigen_vectors = eigen_vectors[:, sorted_index[:
    number_of_pc]]

    #Projecting data over the selected number of
    50 principal components
    principal_component = centred_df.dot(selected_eigen_vectors) #Performing dot product
    principal_component.columns = [f'PC{i+1}' for i in

```

```

range(number_of_pc)]
55  ## Creating Scree plots

fig, axes = plt.subplots(1, 2, figsize=(20, 6))

# Plot the first subplot of variance on the left side
60 axes[0].plot(range(1, len(df_number_pc_var)+1),
df_number_pc_var['variance'], 'ro-', linewidth=2)
axes[0].set_title('Scree Plot for Variance')
axes[0].set_ylabel('Variance')
axes[0].set_xlabel('Number of Principal Component')
65

# Plot the second subplot of cumulative variance on
the right side
axes[1].plot(range(1, len(df_number_pc_var)+1),
70 df_number_pc_var['cumulative_variance'], 'ro-',
linewidth=2)
axes[1].set_title('Scree Plot for Cumulative
Variance')
axes[1].set_ylabel('Cumulative Variance')
75 axes[1].set_xlabel('Number of Principal Component')

plt.show()
print('Correlation of PC1, PC2
\n', principal_component[['PC1', 'PC2']].corr())
80 plt.scatter(principal_component['PC1'], principal_comp
onent['PC2'])
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.title('Scatter plot of PC1 VS PC2')
85 plt.show()
print('The Total variance explained by PC1, PC2 is {}
percent'.format(round(np.real(df_number_pc_var['cumul
ative_variance'][1])*100, 2)))

90 loadings =
pd.DataFrame(selected_eigen_vectors, index=df.columns)
warnings.filterwarnings("ignore")

return
95 principal_component, loadings, df_variance, df_number_pc
_var, eigen_values, eigen_vectors

principal_component, loadings, df_variance, df_number_pc_var
, eigen_values, eigen_vectors = PCA(df_cleaned_dia, 0.9)
100 df_cleaned_dia =
pd.DataFrame(np.real(principal_component.values),
columns=principal_component.columns)
105

```

```

from sklearn.metrics import silhouette_score,
calinski_harabasz_score
def GMM_initialization(df,k):
110     number_of_rows=df.shape[0]
        number_of_columns=df.shape[1]
        means_matrix = df.sample(n=k).values
        identity_matrix=np.eye(number_of_columns)
        covariance_matrix=np.array([identity_matrix]*k)
115     weights_matrix=np.array([float(1/k)]*k)
        return
        means_matrix,covariance_matrix,weights_matrix,number_
        of_rows,number_of_columns

120
def
calculate_posterior(data,means_matrix,covariance_matrix,w
eights_matrix,k,number_of_columns):
    posterior=np.zeros(k)
125     for i in range(k):
        try:
            pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
            np.eye(covariance_matrix[i].shape[0]) * 1e-
130             2,
            rcond=1e-10)
            posterior[i] = multivariate_normal.pdf(data,
            mean=means_matrix[i], cov=pseudo_inverse)
        except Exception as e:
135             continue
    return posterior*weights_matrix/(posterior*weights_matrix).sum()

def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
140     return_counts=True)
    missing_labels=[i for i in range(k) if i not in
    unique_values]
    unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

145     indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
    size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
150         labels[val]=missing_labels[i]
    return labels

def sum_of_square_error_em(new_centroids, data, labels):
155     columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
    # Rename the '0' column of the labels dataframe to
    'Label'

```

```

160     data.rename(columns={0:'Label'}, inplace=True)
        sse = []
        # Compute the distance between each data point and
        its assigned centroid
        for i in range(len(new_centroids)):
165             distance =
                np.sum(np.square(data[data['Label']==i]
                    [columns] - new_centroids.iloc[i]
                    [columns],dtype=np.float64), axis=1)
                #print(distance)
170             sse.append(distance.sum())
        # Return the sum of squared errors

        a=sum(sse)
        return a

175 def Calinski_index_em(df_data,clusters):
        ch_score = calinski_harabasz_score(df_data, clusters)
        return ch_score

180

def GMM(df_cleaned_dia,k,tao):
        scaler = StandardScaler()
        scaler.fit(df_cleaned_dia)
185         scaled_input=scaler.transform(df_cleaned_dia)

        scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

190         means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns=
            GMM_initialization(scaled_input_df,k)
            likelihood=0
            means_matrix_initial=means_matrix
            for i in range(k):
195                 try:

                        pseudo_inverse =
                            np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0]
                                ) * 1e-10))
200                 likelihood=likelihood+weights_matrix[i]*multivariate_normal.logpdf(scaled_input,means_matrix
                    [i], pseudo_inverse)
                except Exception as e:
                    continue

205         log_likelihood_old=np.sum(likelihood)
            old_means_matrix_df=pd.DataFrame(means_matrix)
            posterior_probability = np.zeros((scaled_input.shape[0], k))
            iterations=0

210         while (True):
                iterations+=1
                # Expectation

```

```

215     for i in range(scaled_input.shape[0]):
        posterior_probability[i] =
            calculate_posterior(scaled_input[i],
                                means_matrix, covariance_matrix, weights_matrix,
                                k, number_of_columns)

        # Maximization
220     posterior_probability = np.nan_to_num(posterior_prob
        ability, nan=0)
        for i in range(k):
            # Calculating weight
            weight = posterior_probability[:, i].sum()
225             #print(weight)
            # Updating each centroid
            means_matrix[i] = (posterior_probability[:,
            i] @ scaled_input) / weight
            #print(1, means_matrix[i])
230             # Subtracting the mean value from data
            scaled_input_diff = scaled_input - means_matrix[i]

            # Update the covariance matrix
            covariance_matrix[i] =
235             (posterior_probability[:, i] *
            scaled_input_diff.T @ scaled_input_diff) /
            weight

            # Update the weights matrix
240             weights_matrix[i] = weight / number_of_rows

        likelihood=0
        for i in range(k):
245             try:
                pseudo_inverse =
                    np.linalg.pinv(covariance_matrix[i] +
                    np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
                likelihood=likelihood+weights_matrix[i]*m
250                 ultivariate_normal.logpdf(scaled_input, me
                ans_matrix[i], sudo_inverse)
            except Exception as e:
                continue
        log_likelihood_new = np.sum(likelihood)
255
        new_means_matrix_df = pd.DataFrame(means_matrix)
        distance = []
        for col in new_means_matrix_df.columns:
            col_distance =
260             euclidean(old_means_matrix_df[col],
            new_means_matrix_df[col])
            distance.append(col_distance)
        tao_calculated = sum(distance)/k
265

```

```

270     if tao_calculated<
        tao:#log_likelihood_new>log_likelihood_old and
        100*((log_likelihood_new - log_likelihood_old) /
        log_likelihood_old)<tao:

        print ("Converged")
        labels=np.argmax(posterior_probability,axis=1
        )
275     labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
280     scaled_input_df, labels_df)
        clainski=
        Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial
    #else:
285     #log_likelihood_old=log_likelihood_new

    if iterations>100:
        print ("Max iteration reached")
        labels=np.argmax(posterior_probability,axis=1
290     )
        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
295     sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
        clainski=
        Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial
300

import time
from scipy.spatial.distance import euclidean
def initialize_centroids(df, k,means_matrix):
305     """
        Function to initialize random centroids from dataset.
        Input:
            - df: pandas dataframe with the data
            - k: integer number of clusters
310     Output:
            - temp_df: pandas dataframe with the centroids as columns and index as label
        """

315     centroids=pd.DataFrame(means_matrix,columns=df.column
        s)
        centroids=centroids.T
        centroids.index.name = 'Label'

```



```

320     return centroids

def assign_labels(df, centroids):
    """
325     Function to calculate the closest centroid label for each row in a dataframe.
    Input:
        - df: pandas dataframe with the data
        - centroids: pandas dataframe with the centroids as columns and index as label
    Output:
330        - distances.idxmin(axis=1): pandas series with
            the label of the closest centroid for each row
            in df
    """
    distances = centroids.swifter.apply(lambda x:
335    np.sqrt(((df - x) ** 2).sum(axis=1))) # Calculate
    the Euclidean distance between each row in df and
    each centroid
    return distances.idxmin(axis=1) # Get the index of
340    the minimum distance, which corresponds to the label
    of the closest centroid

def new_centroids(df_label, df1):
    """
345    Function to calculate the new centroids based on the
    current labels of the rows.
    Input:
        - df_label: pandas series with the label of the
            closest centroid for each row in df1
350        - df1: pandas dataframe with the data
    Output:
        - new_centroids.T: pandas dataframe with the new
            centroids as columns and index as feature name
    """
355    joined_df = df1.join(df_label)
    joined_df.rename(columns={0: 'Label'}, inplace=True) # Rename the column with the label
    # Calculate the mean of the rows with the same label
    return joined_df.groupby('Label').mean().T #
    Transpose the dataframe to have the new centroids as
360    columns and index as feature name

def error_clusters(df_new_centroids, df1, df_label):
    """
365    Calculate the error rate of each cluster.

    Args:
        - df_label (pandas.DataFrame): the label of the
370        nearest centroid for each data point.
        - df1 (pandas.DataFrame): the dataset.

```

```

- df_new_centroids (pandas.DataFrame): The new centroids computed in the current iteration.

Returns:
375 - error_rate (float): the total error rate of all clusters.
    """

    #Calculate mean value
380 mean_centroid=df1.groupby('readmitted').mean().reset_index()
    # Transpose the new centroids dataframe and reset the index
    new_centroids= df_new_centroids.T
    # Get the columns of the data dataframe
    columns = df1.columns

385
    sse = []
    # Compute the distance between each data point and its assigned centroid
    for i in range(len(new_centroids)):      ### centroid
        s=[]
390         for j in range(len(mean_centroid)): ### mean centroid
            # Compute the distance between each data point and its assigned centroid
            distance =
                np.sum(np.square(mean_centroid[mean_centroid[
                    'readmitted']==j][columns] -
395                 new_centroids.iloc[i][columns]), axis=1)
                s.append(distance.iloc[0])
            sse.append(s)
        ## key is the cluster number and value is the merged value
        merge_label=pd.DataFrame(sse).idxmin(axis=1).to_dict()
400        ## Merging cluster based on the target variable
        df_label[0]=df_label[0].replace(merge_label)

    df1 = df1.join(df_label) # add the label column to the dataset
    df1.rename(columns={0: 'Label'}, inplace=True) # rename the label column
405    error_list = []
    for i in df1['Label'].value_counts().index:
        df_cluster = df1[df1['Label'] == i] # filter the
            dataset to include only the data points in the
            current cluster
410        y = len(df_cluster[df_cluster['readmitted'] ==
            1]) # count the number of data points in the
            current cluster that were readmitted
            n = len(df_cluster[df_cluster['readmitted'] ==
            0]) # count the number of data points in the
415        current cluster that were not readmitted
            if y == 0 and n == 0:
                error = 0
            else:
                error = n / (n + y) # calculate the error
420                rate of the current cluster
            error_list.append(error)
    return round(sum(error_list),4)

```

```

425 def sum_of_square_error(new_centroids, data, labels):
    """
    Computes the sum of squared errors between the data
    points and their assigned centroids.

430    Args:
    new_centroids (DataFrame): The new centroids computed in the current iteration.
    data (DataFrame): The input data points.
    labels (DataFrame): The labels assigned to each data point.

435    Returns:
    The sum of squared errors.
    """
    # Transpose the new centroids dataframe and reset the index
    new_centroids = new_centroids.T.reset_index()
440    # Get the columns of the data dataframe
    columns = new_centroids.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
    # Rename the '0' column of the labels dataframe to 'Label'
445    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
    its assigned centroid
    for i in range(len(new_centroids)):
450        distance =
            np.sum(np.square(data[data['Label']==i][columns]
                - new_centroids.iloc[i][columns]), axis=1)
            sse.append(sum(distance))
    # Return the sum of squared errors
455    return sum(sse)

def Calinski_index(df_data, clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score
460

def kmeans_lyod_with_error(df1, k, tou, means_matrix_initial):
465    """
    Function to run the K-means Lloyd algorithm.
    Input:
        - df1: pandas dataframe with the data
        - k: integer number of clusters
470        - tou: float tolerance level to stop the algorithm
    Output:
        - centroids: pandas dataframe with the final centroids as columns and index as label
    """
    start_time=time.time()
475    centroids = initialize_centroids(df1,
        k, means_matrix_initial) # Initialize random centroids
    initial_list_of_columns = centroids.columns.to_list()

```

```

iteration = 0
while True:
    # Assign labels to current centroids
    df_label = assign_labels(df1, centroids)
    df_label = pd.DataFrame(df_label)
    # Calculate new centroids
    df_new_centroids = new_centroids(df_label, df1)
    new_list_of_columns =
    df_new_centroids.columns.to_list()
    # Keep the number of clusters the same i.e
    maintain same k
    for i in initial_list_of_columns:
        if i not in new_list_of_columns:
            df_new_centroids[i] = centroids[i]
    # Calculate tao
    distance = []
    for col in centroids.columns:
        col_distance = euclidean(centroids[col], df_new_centroids[col])
        distance.append(col_distance)
    tao_calculated=sum(distance)/k #Used the formula provided for calculating Tao
    sse = sum_of_square_error(df_new_centroids, df1, df_label)
    #error=error_clusters(df_label,df1,k)
    end_time= time.time()
    clainski= Calinski_index(df1,df_label)
    if iteration>100:
        print("Iteration exceeded")

    return sse,clainski
    break

    if tao_calculated<tou or iteration >100:    #if
    the convergence is met, kmeans will stop or
    else if the convergence is never met, after 100
    iteration code will stop
        return sse,clainski
        break                                # otherwise indefinite loop
    else:
        centroids= df_new_centroids # In case we
        need more iterations, the centroids
        calculated at this step acts as input
        iteration+=1

scaler = StandardScaler()
scaler.fit(df_cleaned_dia)
scaled_input=scaler.transform(df_cleaned_dia)

scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

error_matrix_em=[]
error_matrix_kmeans=[]

```

```

for i in range(2,6):
    for j in range(1,21):
        sse,clainski,means_matrix_initial=GMM(df_cleaned_
        dia,i,10)
535     error_matrix_em.append([i,sse,clainski])

        sse,clainski=kmeans_lyod_with_error(scaled_input_
        df,i,10,means_matrix_initial)
        error_matrix_kmeans.append([i,sse,clainski])
540 error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse','clainski'])
error_df_kmeans= pd.DataFrame(error_matrix_kmeans,columns=
['number_of_cluster', 'sse','clainski'])

545 error_df_em.to_csv('6_em_pca.csv',index=False)
error_df_kmeans.to_csv('6_kmeans_pca.csv',index=False)

error_df_em['algo']='em'
550 error_df_kmeans['algo']='kmeans'

run_time_diab=pd.DataFrame()
run_time_diab=pd.concat( [ error_df_em[['algo','number_of_cluster','sse','clainski']],
555     error_df_kmeans[['algo','number_of_cluster', 'sse','clainski']]
        ],ignore_index=True )

import seaborn as sns

560 fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin (['kmeans','em'])],ax=ax);
plt.title('Box Plot of SSE for GMM and K means after PCA')
565 plt.show()

import seaborn as sns

fig, ax = plt.subplots(figsize=(8,6))
570 sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin (['kmeans','em'])],ax=ax);
plt.title('Box Plot of Clainski Index for GMM and K means after PCA')
plt.show()

575 error_df_em_nopca=pd.read_csv('6_em.csv')
error_df_kmeans_nopca=pd.read_csv('6_kmeans.csv')

580 error_df_em_nopca['algo']='em_nopca'
error_df_kmeans_nopca['algo']='kmeans_nopca'

```

```

run_time_diab=pd.concat( [ run_time_diab[['algo','number_of_cluster','sse','clainski']],
585   error_df_em_nopca[['algo','number_of_cluster','sse','clainski']],
   error_df_kmeans_nopca[['algo','number_of_cluster','sse','clainski']]
                           ],ignore_index=True )

import seaborn as sns
590
fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin (['kmeans','em','kmeans_nopca','em_nopca'])])
595 plt.title('Box Plot of SSE for GMM and K means before and after PCA')
plt.show()

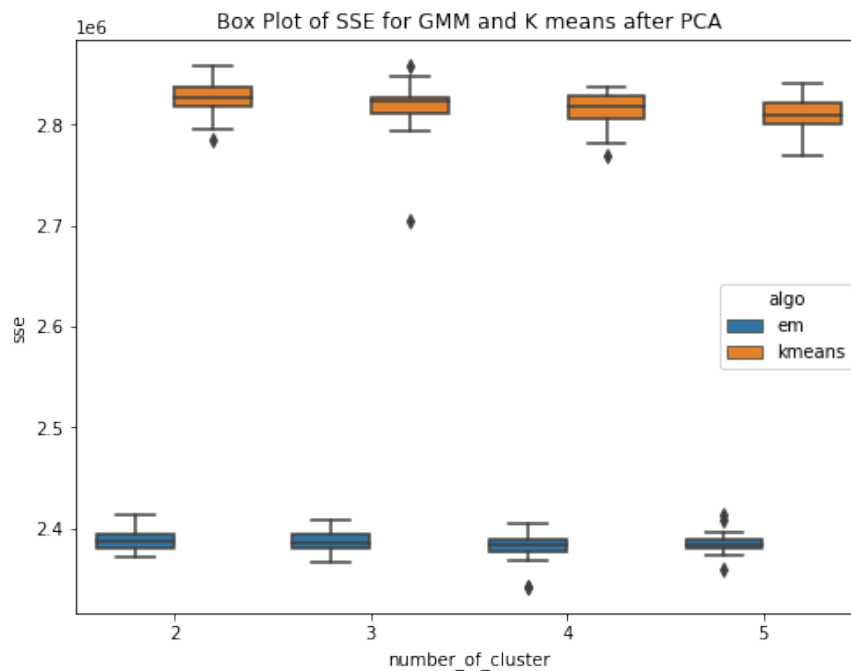
import seaborn as sns
600
fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin (['kmeans','em','kmeans_nopca','em_nopca'])])
605
plt.title('Box Plot of Clainski for GMM and K means before and after PCA')
plt.show()

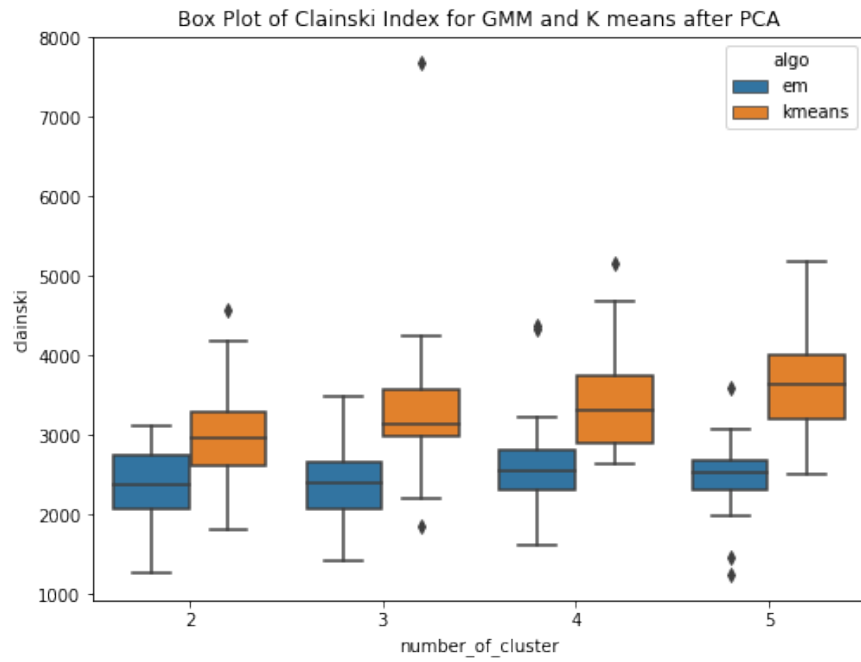
```

## Plot/s

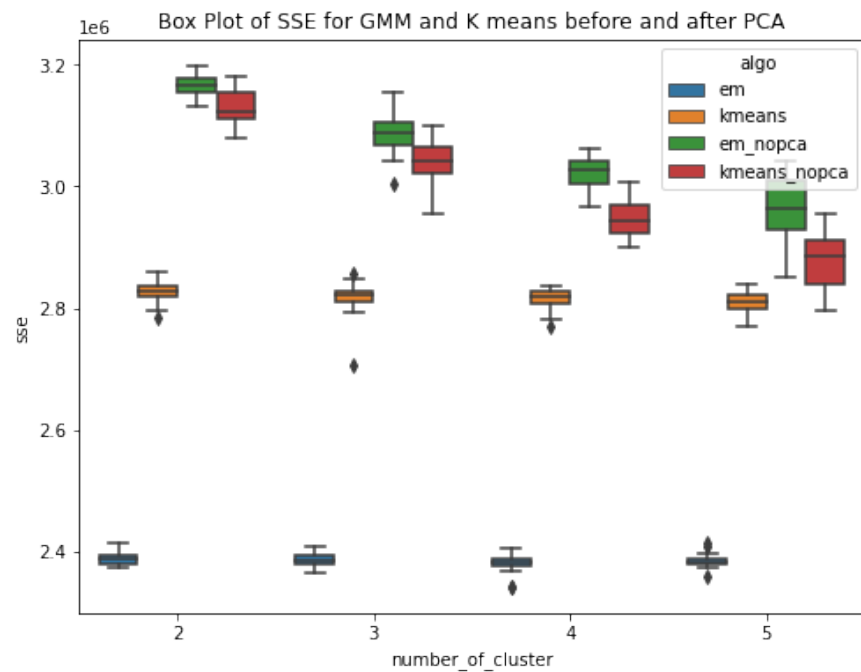
Place images here with suitable captions.



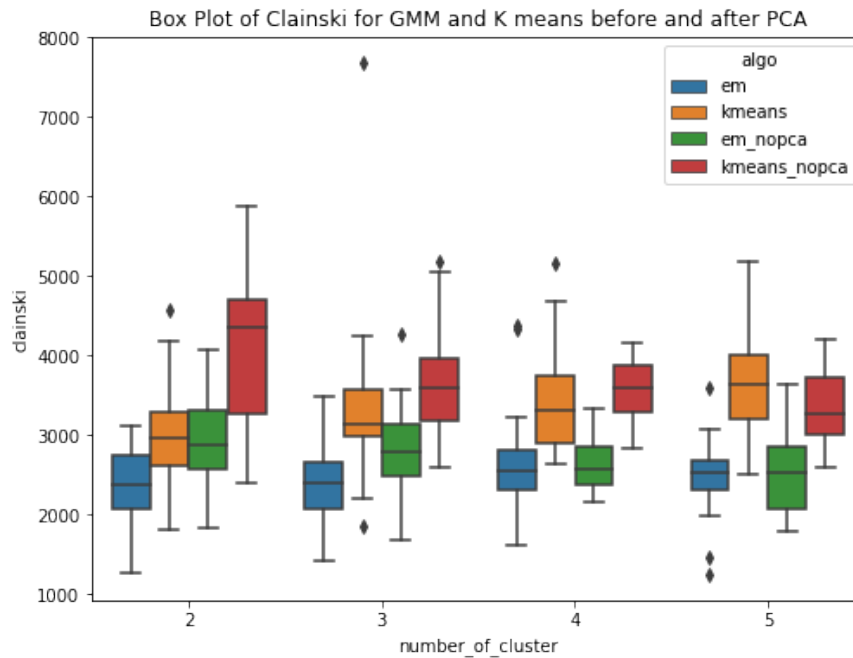
(7)



(8)



(9)



(10)

## Discussion of Experiments

Answer here. . . Have performed PCA over the diabetes dataset. To cover 90 percent of the variance, 24 principal components are required. So the dataset here gets reduced and now we use these principal components as data to run our GMM (EM) and kmeans algorithm. Have ran the algorithms for  $k = 2, 3, 4, 5$  and for each cluster the experiment is performed 20 times. PCA (Principal Component Analysis) is a method used for reducing the dimensionality of a dataset by transforming it into a lower-dimensional space while preserving most of the data variability. The transformed features resulting from PCA are linear combinations of the original features that are orthogonal to each other. PCA can help approximate a normal distribution of the data by removing redundant and correlated features and reducing the effects of outliers and noise. The aim is to improve the results. This is confirmed by the box plots. The box plot showing the within SSE between Kmeans and GMM shows the performance of GMM is much better than that of Kmeans after PCA. The median with sse is less for GMM as compared to that of the kmeans. The second plot shows the CHS score box plot for GMM and kmeans, showing the median value of CHS for kmeans is higher than that of the GMM. The third plot shows the comparison of within SSE for GMM and Kmeans before and after PCA. The plot shows that the sse after performing pca has much less value as compared to the previous experiments when PCA was not performed. The plot shows a huge difference in values of within sse for GMM with and without PCA and for kmeans. The fourth plot is a similar comparison but using CHS score as the evaluating metric. This shows that CHS score is higher before performing pca and reduces after performing PCA.



4. Run the EM algorithm for the other two different mixture models such as, Poisson. Compare all your three  $G_k$ 's using two different appropriate cluster validity techniques, i.e., internal, external or relative indices. Plots are generally a good way to convey complex ideas quickly, i.e., box plots, whisker plots. Discuss your results [30 points].

## R or Python script

```
# Sample R Script With Highlighting
```

```
# Sample Python Script With Highlighting
```

```
import pandas as pd
import swifter
5 import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
10 from scipy.stats import multivariate_normal
from sklearn.metrics import silhouette_score,
calinski_harabasz_score
from scipy.spatial.distance import euclidean

15 df= pd.read_csv('dataset_diabetes/diabetic_data.csv')

from sklearn.metrics import silhouette_score,
calinski_harabasz_score
20 def GMM_initialization(df,k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
    means_matrix = df.sample(n=k).values
    identity_matrix=np.eye(number_of_columns)
25 covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)
    return
    means_matrix,covariance_matrix,weights_matrix,number_
    of_rows,number_of_columns
30

def calculate_posterior(data,means_matrix,covariance_matrix,w
eights_matrix,k,number_of_columns):
    posterior=np.zeros(k)
35

    gamma_arr = np.array([math.gamma(i+1) for i in data])
    for i in range(k):
        pois = (np.exp(-means_matrix[i]) * means_matrix[i] ** data)/gamma_arr
        posterior[i] = weights_matrix[i]*np.prod(pois)+0.0001
40

    return posterior

def maintain_k_clusters(labels,k):
```

```

45     unique_values, value_counts = np.unique(labels,
        return_counts=True)
    missing_labels=[i for i in range(k) if i not in
        unique_values]
    unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

50     indices=[i for i in range(len(labels)) if
        labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
        size=len(missing_labels), replace=False)
    for i, val in enumerate(random_indices):
55         labels[val]=missing_labels[i]
    return labels

def sum_of_square_error_em(new_centroids, data, labels):
60     columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
    # Rename the '0' column of the labels dataframe to
    'Label'
65     data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
    its assigned centroid
    for i in range(len(new_centroids)):
70         distance =
            np.sum(np.square(data[data['Label']==i]
                [columns] - new_centroids.iloc[i]
                [columns],dtype=np.float64), axis=1)
            #print(distance)
75         sse.append(distance.sum())
    # Return the sum of squared errors

    a=sum(sse)
    return a

80 def Calinski_index_em(df_data, clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score

85

def GMM(df_cleaned_dia, k, tao):
    scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
90     scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input, columns=df_cleaned_dia.colu
        mns)

95     means_matrix, covariance_matrix, weights_matrix, number_of_rows, number_of_columns=
        GMM_initialization(scaled_input_df, k)

```

```

likelihood=0
means_matrix_initial=means_matrix
for i in range(k):
    100     try:

        pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0]
            ) * 1e-10))
    105     likelihood=likelihood+weights_matrix[i]*multi
        variate_normal.logpdf(scaled_input, means_matrix
            [i], pseudo_inverse)
    except Exception as e:
        continue

    110 log_likelihood_old=np.sum(likelihood)
        old_means_matrix_df=pd.DataFrame(means_matrix)
        posterior_probability = np.zeros((scaled_input.shape[0], k))
        iterations=0

    115 while (True):
        iterations+=1
        # Expectation
        for i in range(scaled_input.shape[0]):
            posterior_probability[i] =
    120             calculate_posterior(scaled_input[i],
                means_matrix, covariance_matrix, weights_matrix,
                k, number_of_columns)

        # Maximization
    125 posterior_probability=np.nan_to_num(posterior_prob
        ability, nan=0)
        for i in range(k):
            # Calculating weight
            weight = posterior_probability[:, i].sum()
    130             #print(weight)
            # Updating each centroid
            means_matrix[i] = (posterior_probability[:,
                i] @ scaled_input) / weight
            #print(1, means_matrix[i])
    135             # Subtracting the mean value from data
            scaled_input_diff = scaled_input - means_matrix[i]

            # Update the covariance matrix
            covariance_matrix[i] =
    140             (posterior_probability[:, i] *
                scaled_input_diff.T @ scaled_input_diff) /
                weight

            # Update the weights matrix
    145             weights_matrix[i] = weight / number_of_rows

        likelihood=0
        for i in range(k):

```

```

150         try:
            pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
            np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*m
155         multivariate_normal.logpdf(scaled_input,me
            ans_matrix[i], sudo_inverse)
        except Exception as e:
            continue
    log_likelihood_new =np.sum(likelihood)

160
    new_means_matrix_df=pd.DataFrame(means_matrix)
    distance = []
    for col in new_means_matrix_df.columns:
        col_distance =
165         euclidean(old_means_matrix_df[col],
            new_means_matrix_df[col])
        distance.append(col_distance)
    tao_calculated=sum(distance)/k

170

    if tao_calculated<
    tao:#log_likelihood_new>log_likelihood_old and
    100*((log_likelihood_new - log_likelihood_old) /
175    log_likelihood_old)<tao:

        print ("Converged")
        labels=np.argmax(posterior_probability,axis=1
        )
180        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
185        scaled_input_df, labels_df)
        clainski=
        Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial
    #else:
190        #log_likelihood_old=log_likelihood_new

    if iterations>100:
        print ("Max iteration reached")
        labels=np.argmax(posterior_probability,axis=1
195        )
        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
200        scaled_input_df, labels_df)
        clainski=

```

```

        Calinski_index_em(scaled_input_df, labels_df)
        return sse, clainiski, means_matrix_initial
205
error_matrix_em=[]

for i in range(2,6):
210     for j in range(1,21):
        sse, clainiski, means_matrix_initial=GMM(df_cleaned_
            dia, i, 10)
        error_matrix_em.append([i, sse, clainiski])

215 error_df_em= pd.DataFrame(error_matrix_em, columns=['number_of_cluster', 'sse', 'clainiski'])

error_df_em.to_csv('6_em_poisson.csv', index=False)

import pandas as pd
220 import swifter
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
225 from sklearn.preprocessing import StandardScaler
from scipy.stats import multivariate_normal
from sklearn.metrics import silhouette_score,
calinski_harabasz_score
from scipy.spatial.distance import euclidean
230
df= pd.read_csv('dataset_diabetes/diabetic_data.csv')

from sklearn.metrics import silhouette_score,
235 calinski_harabasz_score
def GMM_initialization(df, k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
    means_matrix = df.sample(n=k).values
240    identity_matrix=np.eye(number_of_columns)
    covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)
    return
    means_matrix, covariance_matrix, weights_matrix, number_
245 of_rows, number_of_columns

def calculate_posterior(data, means_matrix, covariance_matrix, w
250 eights_matrix, k, number_of_columns):
    posterior=np.zeros(k)

    for i in range(k):
        m=1/(means_matrix[i]+0.01)
255        exp= np.exp(-m*data)/m

```

```

        posterior[i] = weights_matrix[i]*np.prod(exp)+0.0001
    return posterior

260 def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
    return_counts=True)
    missing_labels=[i for i in range(k) if i not in
    unique_values]
265 unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

    indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
270 size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
        labels[val]=missing_labels[i]
    return labels

275 def sum_of_square_error_em(new_centroids, data, labels):
    columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
280 # Rename the '0' column of the labels dataframe to
    'Label'
    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
285 its assigned centroid
    for i in range(len(new_centroids)):
        distance =
        np.sum(np.square(data[data['Label']==i]
        [columns] - new_centroids.iloc[i]
290 [columns],dtype=np.float64), axis=1)
        #print(distance)
        sse.append(distance.sum())
    # Return the sum of squared errors

295 a=sum(sse)
    return a

def Calinski_index_em(df_data,clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
300 return ch_score

def GMM(df_cleaned_dia,k,tao):
305 scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

```

```

310 scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.colu
    mns)

    means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns=
    GMM_initialization(scaled_input_df,k)
    likelihood=0
315 means_matrix_initial=means_matrix
    for i in range(k):
        try:

            pseudo_inverse =
320 np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0]
                ) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*multi
            variate_normal.logpdf(scaled_input,means_matrix
                [i], pseudo_inverse)
325 except Exception as e:
                continue

    log_likelihood_old=np.sum(likelihood)
    old_means_matrix_df=pd.DataFrame(means_matrix)
    posterior_probability = np.zeros((scaled_input.shape[0], k))
330 iterations=0

    while (True):
        iterations+=1
        # Expectation
335 for i in range(scaled_input.shape[0]):
            posterior_probability[i] =
            calculate_posterior(scaled_input[i],
                means_matrix,covariance_matrix,weights_matrix,
                k,number_of_columns)

340
        # Maximization
        posterior_probability=np.nan_to_num(posterior_prob
            ability, nan=0)
        for i in range(k):
345             # Calculating weight
            weight = posterior_probability[:, i].sum()
            #print(weight)
            # Updating each centroid
            means_matrix[i] = (posterior_probability[:,
350 i] @ scaled_input) / weight
            #print(1,means_matrix[i])
            # Subtracting the mean value from data
            scaled_input_diff = scaled_input - means_matrix[i]

355
            # Update the covariance matrix
            covariance_matrix[i] =
            (posterior_probability[:, i] *
                scaled_input_diff.T @ scaled_input_diff) /
            weight

360
            # Update the weights matrix

```

```

        weights_matrix[i] = weight / number_of_rows

365     likelihood=0
    for i in range(k):
        try:
            pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
370             np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*m
            ultivariate_normal.logpdf(scaled_input,me
            ans_matrix[i], sudo_inverse)
        except Exception as e:
375             continue
    log_likelihood_new =np.sum(likelihood)

    new_means_matrix_df=pd.DataFrame(means_matrix)
    distance = []
380     for col in new_means_matrix_df.columns:
        col_distance =
        euclidean(old_means_matrix_df[col],
        new_means_matrix_df[col])
        distance.append(col_distance)
385     tao_calculated=sum(distance)/k

    if tao_calculated<
390     tao:#log_likelihood_new>log_likelihood_old and
    100*((log_likelihood_new - log_likelihood_old) /
    log_likelihood_old)<tao:

        print ("Converged")
395     labels=np.argmax(posterior_probability,axis=1
    )
    labels=maintain_k_clusters(labels,k)
    labels_df=pd.DataFrame(labels)
    means_matrix_df=pd.DataFrame(means_matrix,col
400     umns=scaled_input_df.columns)
    sse=sum_of_square_error_em(means_matrix_df,
    scaled_input_df, labels_df)
    clainski=
    Calinski_index_em(scaled_input_df,labels_df)
405     return sse,clainski,means_matrix_initial
    #else:
        #log_likelihood_old=log_likelihood_new

    if iterations>100:
410         print ("Max iteration reached")
        labels=np.argmax(posterior_probability,axis=1
        )
        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)

```



```

415         means_matrix_df=pd.DataFrame(means_matrix,columns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
        clainski=
420         Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial

error_matrix_em=[]
425 error_matrix_em=[]

for i in range(2,6):
    for j in range(1,21):
430         sse,clainski,means_matrix_initial=GMM(df_cleaned_dia,i,10)
        error_matrix_em.append([i,sse,clainski])

error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse','clainski'])
435 error_df_em.to_csv('6_em_exp.csv',index=False)

from sklearn.metrics import silhouette_score,
440 calinski_harabasz_score
def GMM_initialization(df,k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
    means_matrix = df.sample(n=k).values
445    identity_matrix=np.eye(number_of_columns)
    covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)
    return
    means_matrix,covariance_matrix,weights_matrix,number_
450    of_rows,number_of_columns

def
calculate_posterior(data,means_matrix,covariance_matrix,w
455 eights_matrix,k,number_of_columns):
    posterior=np.zeros(k)
    for i in range(k):
        try:
            pseudo_inverse =
460            np.linalg.pinv(covariance_matrix[i] +
            np.eye(covariance_matrix[i].shape[0]) * 1e-
            2,
            rcond=1e-10)
            posterior[i] = multivariate_normal.pdf(data,
465            mean=means_matrix[i], cov=pseudo_inverse)
        except Exception as e:
            continue

```

```

    return posterior*weights_matrix/(posterior*weights_matrix).sum()

470 def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
    return_counts=True)
    missing_labels=[i for i in range(k) if i not in
    unique_values]
475 unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

    indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
    random_indices = np.random.choice(indices,
480 size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
        labels[val]=missing_labels[i]
    return labels

485 def sum_of_square_error_em(new_centroids, data, labels):
    columns = data.columns
    # Join the data dataframe and the labels dataframe
    data = data.join(labels)
490 # Rename the '0' column of the labels dataframe to
    'Label'
    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
    # Compute the distance between each data point and
495 its assigned centroid
    for i in range(len(new_centroids)):
        distance =
        np.sum(np.square(data[data['Label']==i]
        [columns] - new_centroids.iloc[i]
500 [columns],dtype=np.float64), axis=1)
        #print(distance)
        sse.append(distance.sum())
    # Return the sum of squared errors

505 a=sum(sse)
    return a

def Calinski_index_em(df_data,clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
510 return ch_score

def GMM(df_cleaned_dia,k,tao):
515 scaler = MinMaxScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.colu
520 mns)

```

```

means_matrix, covariance_matrix, weights_matrix, number_of_rows, number_of_columns=
GMM_initialization(scaled_input_df, k)
likelihood=0
525 means_matrix_initial=means_matrix
for i in range(k):
    try:

        pseudo_inverse =
530 np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0]
        ) * 1e-10))
        likelihood=likelihood+weights_matrix[i]*multi
        variate_normal.logpdf(scaled_input, means_matrix
        [i], pseudo_inverse)
535 except Exception as e:
        continue
log_likelihood_old=np.sum(likelihood)
old_means_matrix_df=pd.DataFrame(means_matrix)
posterior_probability = np.zeros((scaled_input.shape[0], k))
540 iterations=0

while (True):
    iterations+=1
    # Expectation
545 for i in range(scaled_input.shape[0]):
        posterior_probability[i] =
        calculate_posterior(scaled_input[i],
        means_matrix, covariance_matrix, weights_matrix,
        k, number_of_columns)

550 # Maximization
posterior_probability=np.nan_to_num(posterior_prob
ability, nan=0)
for i in range(k):
555 # Calculating weight
weight = posterior_probability[:, i].sum()
#print(weight)
# Updating each centroid
means_matrix[i] = (posterior_probability[:,
560 i] @ scaled_input) / weight
#print(1, means_matrix[i])
# Subtracting the mean value from data
scaled_input_diff = scaled_input - means_matrix[i]

565 # Update the covariance matrix
covariance_matrix[i] =
(posterior_probability[:, i] *
scaled_input_diff.T @ scaled_input_diff) /
weight

570 # Update the weights matrix
weights_matrix[i] = weight / number_of_rows

```

```

575     likelihood=0
    for i in range(k):
        try:
            pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] +
580             np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*m
            multivariate_normal.logpdf(scaled_input,me
            ans_matrix[i], sudo_inverse)
        except Exception as e:
585             continue
    log_likelihood_new =np.sum(likelihood)

    new_means_matrix_df=pd.DataFrame(means_matrix)
    distance = []
590    for col in new_means_matrix_df.columns:
        col_distance =
        euclidean(old_means_matrix_df[col],
        new_means_matrix_df[col])
        distance.append(col_distance)
595    tao_calculated=sum(distance)/k

    if tao_calculated<
600    tao:#log_likelihood_new>log_likelihood_old and
    100*((log_likelihood_new - log_likelihood_old) /
    log_likelihood_old)<tao:

        print ("Converged")
605        labels=np.argmax(posterior_probability,axis=1
        )
        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
610        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
        clainski=
        Calinski_index_em(scaled_input_df,labels_df)
615        return sse,clainski,means_matrix_initial
    #else:
        #log_likelihood_old=log_likelihood_new

    if iterations>100:
620        print ("Max iteration reached")
        labels=np.argmax(posterior_probability,axis=1
        )
        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
625        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)

```

```

        sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
        clainski=
630     Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial

error_df_em.to_csv('6_em_minmax.csv',index=False)
635
error_df_em_poisson=pd.read_csv('6_em_poisson.csv')
error_df_em_exp=pd.read_csv('6_em_exp.csv')
error_df_em_normal=pd.read_csv('6_em_minmax.csv')

640 error_df_em_poisson['algo']='em_poisson'
error_df_em_exp['algo']='em_exponential'
error_df_em_normal['algo']='em_normal'

run_time_diab=pd.DataFrame()
645 run_time_diab=pd.concat( [ error_df_em_poisson[['algo','number_of_cluster','sse','clainski']],
        error_df_em_exp[['algo','number_of_cluster','sse','clainski']],
        error_df_em_normal[['algo','number_of_cluster','sse','clainski']]
        ],ignore_index=True )

650 import seaborn as sns

fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
655         data=run_time_diab[run_time_diab['algo'].isin(['em_poisson','em_exponential','em_no
plt.title('Box Plot of SSE for 3 versions of GMM (Normal, Poisson, Exponential)')
plt.show()

660 import seaborn as sns

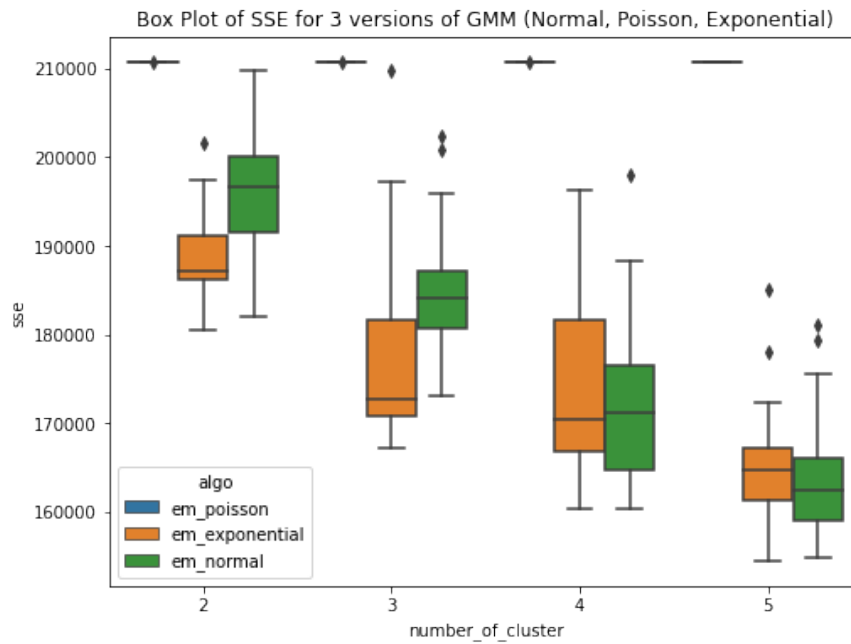
fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
665         data=run_time_diab[run_time_diab['algo'].isin(['em_poisson','em_exponential','em_no
plt.title('Box Plot of Clainski for 3 versions of GMM (Normal, Poisson, Exponential)')
plt.show()

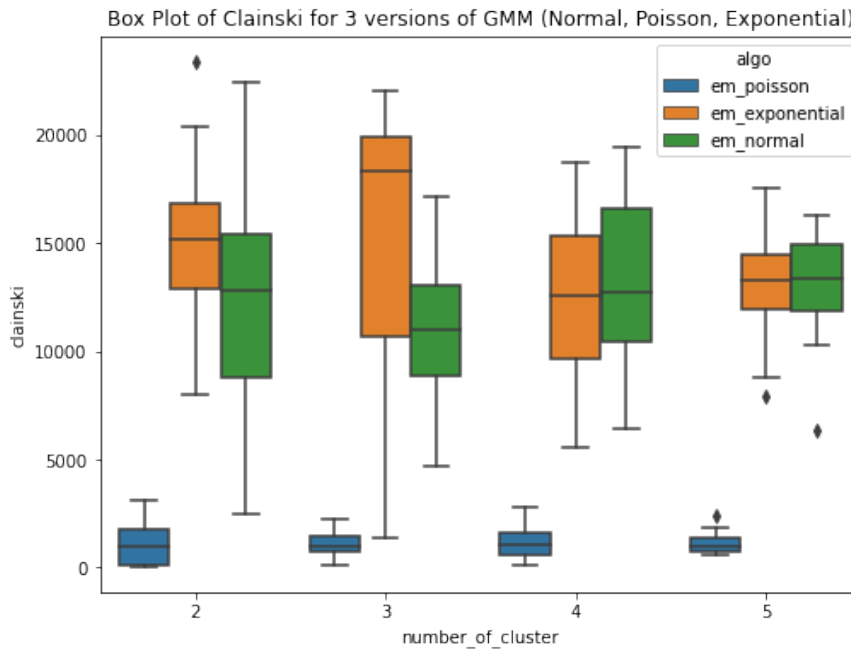
```

## Plot/s

Place images here with suitable captions.



(11)



(12)

## Discussion of Experiments

Answer here... In this question, other than normal distribution, I have tried the following distributions:

1) Poisson Distribution

2) Exponential Distribution Here the Poisson distribution requires positive values for data, so have used MinMaxScaler to scale the data and then have created 3 different codes, one for each distribution type i.e normal, Poisson and exponential. Equation for Poisson Distribution  $P(X=k) = \frac{\lambda^k e^{-\lambda}}{k!}$

Equation for Exponential Distribution  $f(x; \lambda) = \lambda e^{-\lambda x}$

From the box plot of within sum of square error, it can be seen that with increase in the number of clusters, the within SSE of normal distribution improves and for  $k=5$  it is the best. The within sse for poisson is highest. The reason could be the data is normally distributed and hence the assumption that data is normally distributed is correct. Also, in previous assignments EDA, it was identified the data was normally distributed and was skewed. Similarly, the CHS score shows that with increase in number of cluster, the score improves and for  $k=4$  and  $5$  the normal distribution outperforms other distributions (poisson and exponential).

## Problem 3

Improve the EM algorithm through initialization. *k-means++* is an extended *k*-means clustering algorithm and induces non-uniform distributions over the data that serve as the initial centroids. Read the paper and implement this idea to improve your  $G_k$  program. Let's call the new algorithm  $G_{k++}$ . Run your new  $G_{k++}$  and  $G_k$  for  $k = 2, \dots, 5$  for 20 runs each. Compare  $G_k$  and  $C_k$  using two different appropriate cluster validity techniques, i.e., internal, external or relative indices. Plots are generally a good way to convey complex ideas quickly, i.e., box plots, whisker plots. Discuss your results.

### R or Python script

```
# Sample R Script With Highlighting
```

```
# Sample Python Script With Highlighting
import pandas as pd
import swifter
import matplotlib.pyplot as plt
5 import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from tqdm import tqdm
from sklearn.preprocessing import StandardScaler
from scipy.stats import multivariate_normal
10 from sklearn.metrics import silhouette_score, calinski_harabasz_score
from scipy.spatial.distance import euclidean

df= pd.read_csv('dataset_diabetes/diabetic_data.csv')

15 def initialize_centroids_plus(df,k):
    """
    Function to calculate the random centroid using
    kmeans ++ technique.
    Input:
    20     - df: pandas dataframe with the data
        - k: number of clusters
    Output:
        - centroid.T: pandas dataframe with all
          centroids initialized
    """
    25
```

```

#Initialize random centroids from dataset
centroid = []
centroid.append(df.apply(lambda x:
float(x.sample()))))
30 centroid=pd.DataFrame(centroid)
column=centroid.columns.to_list()
##Randomly created first centroid from the domain of
each column in the dataframe

35 for i in range(1,k):
## The above for loop is for generating k-1 clusters
as first random cluster is already generated
    distance=pd.DataFrame()
## Creating dataframe of distance. This will store
the distance of each datapoint from each cluster
    for j in range(len(centroid)):
##This for loop is for finding distance from each
centroid
        a=pd.DataFrame([np.sqrt(np.sum(np.square(df[c
45 column] - centroid.iloc[j][column]),
axis=1))]).T

        distance=pd.concat([distance,a],axis=1)
## Distance_min stores the minimum distance of each point from all the centroids
50 distance_min=distance.min(axis=1)
##Calculates probability for each datapoint
    probability = distance_min / distance_min.sum()
## Selecting the next centroid based on the
probability which is proportional to find square
distance
55 new_centroid =
    pd.DataFrame(df.iloc[np.random.choice(len(df),p=p
robability)]).T

60 centroid=pd.concat([centroid,new_centroid],ignore
_index=True)
    centroid.index.name='Label'
## Concatenated the centroid dataframe with new
centroid and loop continues until all k centroids
65 are initialized randomly
    return centroid

from sklearn.metrics import silhouette_score,
70 calinski_harabasz_score
def GMM_initialization(df,k):
    number_of_rows=df.shape[0]
    number_of_columns=df.shape[1]
    a = initialize_centroids_plus(df,k)
75 means_matrix= a.to_numpy()
    identity_matrix=np.eye(number_of_columns)
    covariance_matrix=np.array([identity_matrix]*k)
    weights_matrix=np.array([float(1/k)]*k)

```



```

    return
80     means_matrix, covariance_matrix, weights_matrix, number_
        of_rows, number_of_columns

def
85 calculate_posterior(data, means_matrix, covariance_matrix, w
    eights_matrix, k, number_of_columns):
    posterior=np.zeros(k)
    for i in range(k):
        try:
90             pseudo_inverse =
                np.linalg.pinv(covariance_matrix[i] +
                np.eye(covariance_matrix[i].shape[0]) * 1e-
                2,
                rcond=1e-10)
95             posterior[i] = multivariate_normal.pdf(data,
                mean=means_matrix[i], cov=pseudo_inverse)
        except Exception as e:
            continue
    return posterior*weights_matrix/(posterior*weights_matrix).sum()
100

def maintain_k_clusters(labels,k):
    unique_values, value_counts = np.unique(labels,
    return_counts=True)
    missing_labels=[i for i in range(k) if i not in
105 unique_values]
    unique_values_to_be_replaced=[unique_values[i] for i in np.where(value_counts > 1)[0]]

    indices=[i for i in range(len(labels)) if
    labels[i] in unique_values_to_be_replaced]
110 random_indices = np.random.choice(indices,
    size=len(missing_labels), replace=False)
    for i,val in enumerate(random_indices):
        labels[val]=missing_labels[i]
    return labels
115

def sum_of_square_error_em(new_centroids, data, labels):
    columns = data.columns
    # Join the data dataframe and the labels dataframe
120 data = data.join(labels)
    # Rename the '0' column of the labels dataframe to
    'Label'
    data.rename(columns={0:'Label'}, inplace=True)
    sse = []
125 # Compute the distance between each data point and
    its assigned centroid
    for i in range(len(new_centroids)):
        distance =
130         np.sum(np.square(data[data['Label']==i]
            [columns] - new_centroids.iloc[i]
            [columns],dtype=np.float64), axis=1)

```

```

        #print(distance)
        sse.append(distance.sum())
        # Return the sum of squared errors
135
    a=sum(sse)
    return a

def Calinski_index_em(df_data,clusters):
140
    ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score

145 def GMM(df_cleaned_dia,k,tao):
    scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

150
    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)

    means_matrix,covariance_matrix,weights_matrix,number_of_rows,number_of_columns=
    GMM_initialization(scaled_input_df,k)
155
    likelihood=0
    means_matrix_initial=means_matrix
    for i in range(k):
        try:

160
            pseudo_inverse =
            np.linalg.pinv(covariance_matrix[i] + np.diag(np.ones(covariance_matrix[i].shape[0])
            ) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*multi
            variate_normal.logpdf(scaled_input,means_matrix
165
            [i], pseudo_inverse)
        except Exception as e:
            continue
    log_likelihood_old=np.sum(likelihood)
    old_means_matrix_df=pd.DataFrame(means_matrix)
170
    posterior_probability = np.zeros((scaled_input.shape[0], k))
    iterations=0

    while (True):
        iterations+=1
175
        # Expectation
        for i in range(scaled_input.shape[0]):
            posterior_probability[i] =
            calculate_posterior(scaled_input[i],
            means_matrix,covariance_matrix,weights_matrix,
180
            k,number_of_columns)

        # Maximization
        posterior_probability=np.nan_to_num(posterior_probability, nan=0)

```

```

185     for i in range(k):
        # Calculating weight
        weight = posterior_probability[:, i].sum()
        #print(weight)
        # Updating each centroid
190     means_matrix[i] = (posterior_probability[:,
        i] @ scaled_input) / weight
        #print(1,means_matrix[i])
        # Subtracting the mean value from data
        scaled_input_diff = scaled_input - means_matrix[i]
195
        # Update the covariance matrix
        covariance_matrix[i] =
        (posterior_probability[:, i] *
        scaled_input_diff.T @ scaled_input_diff) /
200     weight

        # Update the weights matrix
        weights_matrix[i] = weight / number_of_rows

205
    likelihood=0
    for i in range(k):
        try:
            pseudo_inverse =
210             np.linalg.pinv(covariance_matrix[i] +
            np.diag(np.ones(covariance_matrix[i].shape[0]) * 1e-10))
            likelihood=likelihood+weights_matrix[i]*m
            ultivariate_normal.logpdf(scaled_input,me
            ans_matrix[i], sudo_inverse)
215         except Exception as e:
            continue
    log_likelihood_new =np.sum(likelihood)

    new_means_matrix_df=pd.DataFrame(means_matrix)
220    distance = []
    for col in new_means_matrix_df.columns:
        col_distance =
        euclidean(old_means_matrix_df[col],
        new_means_matrix_df[col])
225        distance.append(col_distance)
    tao_calculated=sum(distance)/k

230
    if tao_calculated<
    tao:#log_likelihood_new>log_likelihood_old and
    100*((log_likelihood_new - log_likelihood_old) /
    log_likelihood_old)<tao:

235
        print ("Converged")
        labels=np.argmax(posterior_probability,axis=1
        )

```

```

        labels=maintain_k_clusters(labels,k)
        labels_df=pd.DataFrame(labels)
240     means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
        clainski=
245     Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial
    #else:
        #log_likelihood_old=log_likelihood_new

250     if iterations>100:
        print("Max iteration reached")
        labels=np.argmax(posterior_probability,axis=1
        )
        labels=maintain_k_clusters(labels,k)
255     labels_df=pd.DataFrame(labels)
        means_matrix_df=pd.DataFrame(means_matrix,col
        umns=scaled_input_df.columns)
        sse=sum_of_square_error_em(means_matrix_df,
        scaled_input_df, labels_df)
260     clainski=
        Calinski_index_em(scaled_input_df,labels_df)
        return sse,clainski,means_matrix_initial

265 import time
from scipy.spatial.distance import euclidean
def initialize_centroids(df, k,means_matrix):
    """
    Function to initialize random centroids from dataset.
270     Input:
        - df: pandas dataframe with the data
        - k: integer number of clusters
    Output:
        - temp_df: pandas dataframe with the centroids as columns and index as label
275     """

    centroids=pd.DataFrame(means_matrix,columns=df.columns
    s)
    centroids=centroids.T
280     centroids.index.name = 'Label'
    return centroids

285
def assign_labels(df, centroids):
    """
    Function to calculate the closest centroid label for each row in a dataframe.
    Input:
290     - df: pandas dataframe with the data

```

```

    - centroids: pandas dataframe with the centroids as columns and index as label
Output:
    - distances.idxmin(axis=1): pandas series with
      the label of the closest centroid for each row
295     in df
    """
distances = centroids.swifter.apply(lambda x:
np.sqrt(((df - x) ** 2).sum(axis=1))) # Calculate
the Euclidean distance between each row in df and
300     each centroid
return distances.idxmin(axis=1) # Get the index of
the minimum distance, which corresponds to the label
of the closest centroid

305 def new_centroids(df_label, df1):
    """
    Function to calculate the new centroids based on the
    current labels of the rows.
310     Input:
        - df_label: pandas series with the label of the
          closest centroid for each row in df1
        - df1: pandas dataframe with the data
    Output:
315         - new_centroids.T: pandas dataframe with the new
          centroids as columns and index as feature name
    """
    joined_df = df1.join(df_label)
    joined_df.rename(columns={0: 'Label'}, inplace=True) # Rename the column with the label
320     # Calculate the mean of the rows with the same label
    return joined_df.groupby('Label').mean().T #
    Transpose the dataframe to have the new centroids as
    columns and index as feature name

325

def error_clusters(df_new_centroids, df1, df_label):
    """
    Calculate the error rate of each cluster.
330
    Args:
        - df_label (pandas.DataFrame): the label of the
          nearest centroid for each data point.
        - df1 (pandas.DataFrame): the dataset.
335        - df_new_centroids (pandas.DataFrame): The new centroids computed in the current iteration.

    Returns:
        - error_rate (float): the total error rate of all clusters.
    """
340
    #Calculate mean value
    mean_centroid=df1.groupby('readmitted').mean().reset_index()

```

```

345     # Transpose the new centroids dataframe and reset the index
new_centroids= df_new_centroids.T
    # Get the columns of the data dataframe
columns = df1.columns

    sse = []
350    # Compute the distance between each data point and its assigned centroid
    for i in range(len(new_centroids)):      #### centroid
        s=[]
        for j in range(len(mean_centroid)): ### mean centroid
            # Compute the distance between each data point and its assigned centroid
355            distance =
                np.sum(np.square(mean_centroid[mean_centroid[
                    'readmitted']==j][columns] -
                    new_centroids.iloc[i][columns]), axis=1)
                s.append(distance.iloc[0])
360        sse.append(s)
        ## key is the cluster number and value is the merged value
        merge_label=pd.DataFrame(sse).idxmin(axis=1).to_dict()
        ## Merging cluster based on the target variable
        df_label[0]=df_label[0].replace(merge_label)
365
        df1 = df1.join(df_label) # add the label column to the dataset
        df1.rename(columns={0: 'Label'}, inplace=True) # rename the label column
        error_list = []
        for i in df1['Label'].value_counts().index:
370            df_cluster = df1[df1['Label'] == i] # filter the
                dataset to include only the data points in the
                current cluster
                y = len(df_cluster[df_cluster['readmitted'] ==
375                1]) # count the number of data points in the
                current cluster that were readmitted
                n = len(df_cluster[df_cluster['readmitted'] ==
                0]) # count the number of data points in the
                current cluster that were not readmitted
                if y == 0 and n == 0:
380                    error = 0
                else:
                    error = n / (n + y) # calculate the error
                    rate of the current cluster
                    error_list.append(error)
385    return round(sum(error_list),4)

def sum_of_square_error(new_centroids, data, labels):
    """
390    Computes the sum of squared errors between the data
        points and their assigned centroids.

        Args:
        new_centroids (DataFrame): The new centroids computed in the current iteration.
395        data (DataFrame): The input data points.
        labels (DataFrame): The labels assigned to each data point.

```

```

Returns:
The sum of squared errors.
"""
# Transpose the new centroids dataframe and reset the index
new_centroids = new_centroids.T.reset_index()
# Get the columns of the data dataframe
columns = new_centroids.columns
# Join the data dataframe and the labels dataframe
data = data.join(labels)
# Rename the '0' column of the labels dataframe to 'Label'
data.rename(columns={0:'Label'}, inplace=True)
sse = []
# Compute the distance between each data point and
its assigned centroid
for i in range(len(new_centroids)):
    distance =
    np.sum(np.square(data[data['Label']==i][columns]
    - new_centroids.iloc[i][columns]), axis=1)
    sse.append(sum(distance))
# Return the sum of squared errors
return sum(sse)

def Calinski_index(df_data, clusters):
    ch_score = calinski_harabasz_score(df_data, clusters)
    return ch_score

def kmeans_lyod_with_error(df1, k, tou, means_matrix_initial):
    """
    Function to run the K-means Lloyd algorithm.
    Input:
    - df1: pandas dataframe with the data
    - k: integer number of clusters
    - tou: float tolerance level to stop the algorithm
    Output:
    - centroids: pandas dataframe with the final centroids as columns and index as label
    """
    start_time=time.time()
    centroids = initialize_centroids(df1,
    k, means_matrix_initial) # Initialize random centroids
    initial_list_of_columns = centroids.columns.to_list()
    iteration = 0
    while True:
        # Assign labels to current centroids
        df_label = assign_labels(df1, centroids)
        df_label = pd.DataFrame(df_label)
        # Calculate new centroids
        df_new_centroids = new_centroids(df_label, df1)
        new_list_of_columns =
        df_new_centroids.columns.to_list()

```

```

450     # Keep the number of clusters the same i.e
    maintain same k
    for i in initial_list_of_columns:
        if i not in new_list_of_columns:
            df_new_centroids[i] = centroids[i]
455     # Calculate tao
    distance = []
    for col in centroids.columns:
        col_distance = euclidean(centroids[col], df_new_centroids[col])
        distance.append(col_distance)
460     tao_calculated=sum(distance)/k #Used the formula provided for calculating Tao
    sse = sum_of_square_error(df_new_centroids, df1, df_label)
    #error=error_clusters(df_label,df1,k)
    end_time= time.time()
    clainski= Calinski_index(df1,df_label)
465     if iteration>100:
        print("Iteration exceeded")

        return sse,clainski
        break
470
    if tao_calculated<tou or iteration >100:    #if
    the convergence is met, kmeans will stop or
    else if the convergence is never met, after 100
    iteration code will stop
475         return sse,clainski
        break                                # otherwise indefinite loop
    else:
        centroids= df_new_centroids # In case we
        need more iterations, the centroids
480         calculated at this step acts as input
        iteration+=1

485     scaler = StandardScaler()
    scaler.fit(df_cleaned_dia)
    scaled_input=scaler.transform(df_cleaned_dia)

    scaled_input_df= pd.DataFrame(scaled_input,columns=df_cleaned_dia.columns)
490

    error_matrix_em=[]
    error_matrix_kmeans=[]
    for i in range(2,6):
495         for j in range(1,21):
            sse,clainski,means_matrix_initial=GMM(df_cleaned_dia,i,10)
            error_matrix_em.append([i,sse,clainski])

            sse,clainski=kmeans_lyod_with_error(scaled_input_df,i,10,means_matrix_initial)
500            error_matrix_kmeans.append([i,sse,clainski])
    error_df_em= pd.DataFrame(error_matrix_em,columns=['number_of_cluster', 'sse', 'clainski'])
    error_df_kmeans= pd.DataFrame(error_matrix_kmeans,columns=['number_of_cluster', 'sse', 'clainski'])

```



```

error_df_em.to_csv('6_em++.csv',index=False)
505 error_df_kmeans.to_csv('6_kmeans++.csv',index=False)

error_df_em.to_csv('6_em.csv',index=False)
error_df_kmeans.to_csv('6_kmeans.csv',index=False)

510 error_df_em_normal=pd.read_csv('6_em.csv')
error_df_kmeans_normal=pd.read_csv('6_kmeans.csv')

error_df_em['algo']='em++'
515 error_df_kmeans['algo']='kmeans++'
error_df_em_normal['algo']='em'
error_df_kmeans_normal['algo']='kmeans'

run_time_diab=pd.DataFrame()
520 run_time_diab=pd.concat( [ error_df_em[['algo','number_of_cluster','sse','clainski']],
    error_df_kmeans[['algo','number_of_cluster','sse','clainski']],
    error_df_em_normal[['algo','number_of_cluster','sse','clainski']],
    error_df_kmeans_normal[['algo','number_of_cluster','sse','clainski']]

525
                                ],ignore_index=True )

import seaborn as sns

530 fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin(['em++','kmeans++'])],ax=ax);
plt.title('Box Plot of SSE for GMM++ and K means++ initialization')
535 plt.show()

import seaborn as sns

fig, ax = plt.subplots(figsize=(8,6))

540 sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin(['em++','kmeans++'])],ax=ax);
plt.title('Box Plot of Clainski for GMM++ and K means++ initialization')
plt.show()

545 import seaborn as sns

fig, ax = plt.subplots(figsize=(8,6))

550 sns.boxplot(x='number_of_cluster', y='sse', hue='algo',
            data=run_time_diab[run_time_diab['algo'].isin(['kmeans','em','em++','kmeans++'])],ax=ax);
plt.title('Box Plot of SSE for GMM and K means normal and with ++ initialization')
plt.show()

555 import seaborn as sns

```

```

fig, ax = plt.subplots(figsize=(8,6))

sns.boxplot(x='number_of_cluster', y='clainski', hue='algo',
560         data=run_time_diab[run_time_diab['algo'].isin(['kmeans','em','em++','kmeans++'])],ax=ax)
plt.title('Box Plot of SSE for GMM and K means normal and with ++ initialization')
plt.show()

```

## Discussion of Experiments

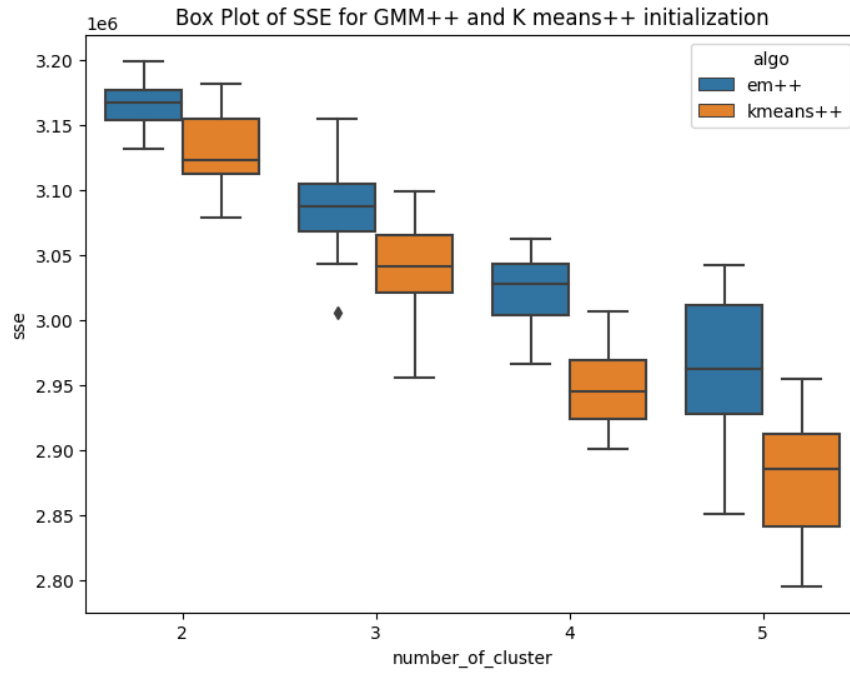
Answer here...

I have implemented the kmeans plus plus algorithm for number of clusters ranging from 2 to 5. The method is similar to kmeans llyod with different centroid initialization. Here the first centroid is chosen randomly from the domain of data. For rest k-1 centroids, we follow a different iterative approach. The second cluster is initialized based on the first centroid. The third centroid requires data of first two centroids. The methodology includes, selecting the first centroid randomly from domain of data. Then second cluster is at the farthest distance from that cluster. In this way iteratively we initialize k centroids. As the centroid initialization is not random, it is expected that the inter cluster distance will be more and intra cluster distance bewteen points and centroid will be less. The K-means++ algorithm selects the centroids in following way Step1: Choosing the first centroid at random from the data points. Step2: For each remaining data point, computing its distance to the nearest centroid that has already been chosen. Step3: Selecting the next centroid randomly from the remaining data points, with probability proportional to the squared distance to the nearest centroid. Repeating steps 2-3 until all K centroids have been chosen. The main idea behind K-means++ initialization is to select centroids that are well spread out across the data points. By selecting the next centroid from the remaining data points with a probability that is proportional to the squared distance to the nearest centroid, K-means++ initialization ensures that data points that are far away from existing centroids are more likely to be selected as new centroids. Hence the sum of square error is expected to be less. Rest the stopping conditions and other steps are similar to that of kmeans llyods.

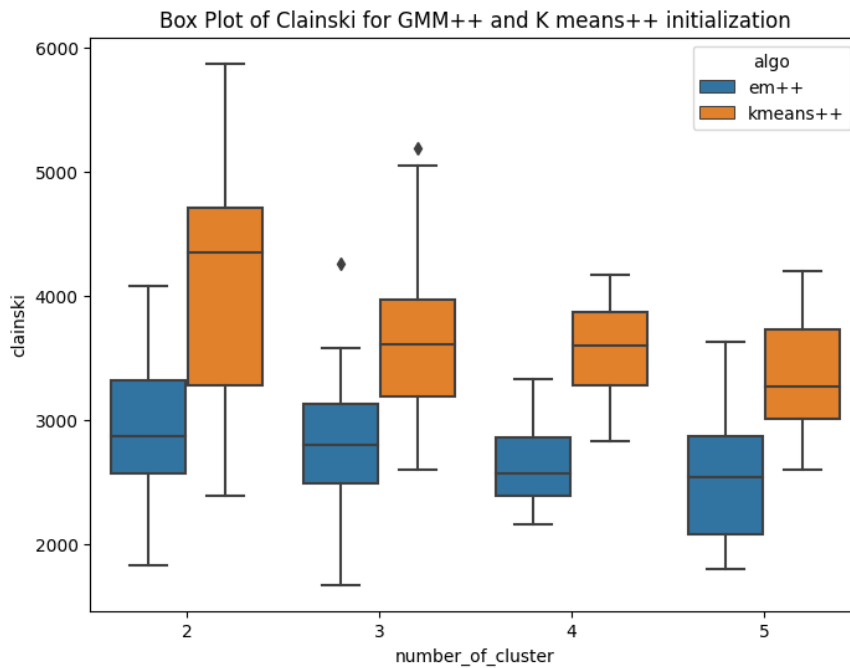
Similarly the mean matrix is initialized for GMM(EM) and same is used for K means. Have ran this experiment for k=2,3,4,5 , running 20 times for each cluster. The first box plot shows the within Sum of square error for GMM plus plus and kmeans plus plus. It can be seen that, the median within SSE of kmeans is less than that of GMM. The second plot shows the CHS score using the new initialization technique and the value of CHS score for kmeans is higher than that of GMM. The third and fourth box plot shows the comparision of sse and CHS score for normal GMM and kmeans with the GMM plus Plus and Kmeans plus plus. The box plots shows that the median error for GMM++ is less than that of normal GMM. And similar observation is achieved for Kmeans algorithm. So by changing the initialization, the performance of the algorithms improved. This is confirmed by the CHS score as well, the median CHS score for plus plus algorithms(GMM and means) is higher than that of the median CHS score for normal initialization of GMM and kmeans

## Plot/s


Place images here with suitable captions.



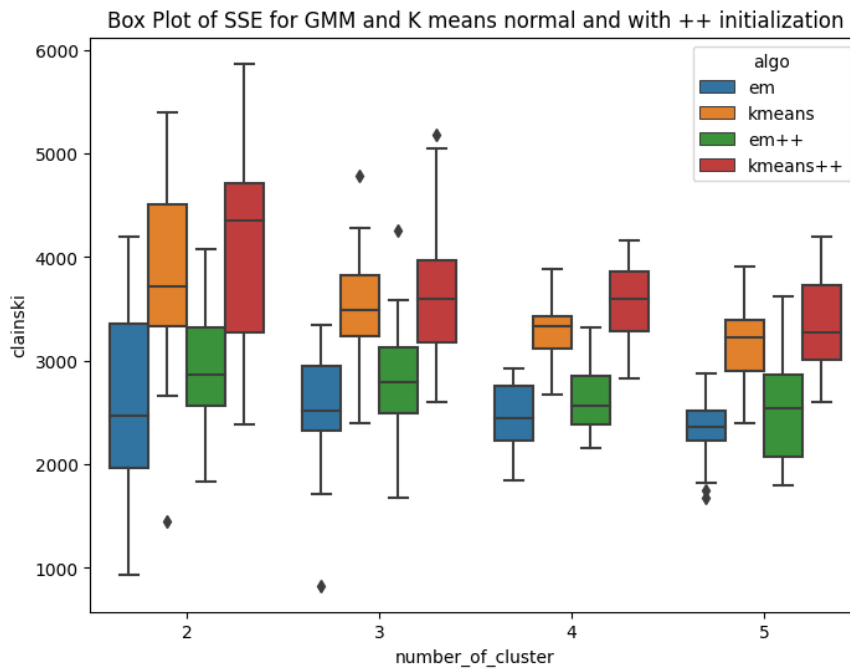
(13)



(14)



(15)



(16)

## Submission

You must use  $\text{\LaTeX}$  to turn in your assignments. Please submit the following two files via Canvas:

1. A .pdf with the name `yourname-hw6-everything.pdf` which you will get after compiling your .tex file.
2. A .zip file with the name `yourname-hw6.zip` which should contain your .tex, .pdf, codes(.py, .ipynb, .R, or .Rmd), and a README file. The README file should contain information about dependencies and how to run your codes.